TOWARD COMPULSORY SCENE-MEDIATED LANGUAGE MODELING: MEANINGFUL SUBSTRATES FOR EVERY NEXT WORD

A PREPRINT

Henrik Westerberg
Emergent Wisdom
henrik.westerberg@emergentwisdom.org

August 12, 2026
Revised: August 25, 2026

ABSTRACT

Can a language model build a world as it writes? Ordinary language models integrate prior words in contextual hidden states and can acquire latent representations of entities and situations, but no externally structured organization is required at their lexical interface. The *Substrate-Translated Language Model* (STLM) makes structured images, descriptions, or videos compulsory for every lexical decision and asks whether this biases computation toward reusable physical and schematic factors. Its development path moves from one fixed substrate per word, through varied context-sensitive instances and short clips, to substrates for longer spans and ultimately one continuously changing scene from which language is read. We call causally useful joint organization that supports composition, manipulation, and reuse across context *scene emergence*; accumulation in one persistent state is its strongest form. This is the long-term hypothesis, not a result demonstrated here.

The stage-1 pilot tests this no-bypass prerequisite on a 1,024-token TinyStories vocabulary with one fixed pixel image or encoded description per token. Predictors emit substrate states from which stateless or causal-sequence readouts predict the next token. On the same 1,024 held-out contexts, causal description and pixel readouts reach 41.9% and 37.8% top-1 accuracy, below a direct-token model at 44.7% but above stateless counterparts at 35.0% and 27.8% and a context-free frequency baseline at 8.1%. The causal configurations differ in architecture, capacity, and supervision density, so their gains do not isolate history. Because each canonical can serve as an arbitrary codeword, the pilot establishes interface feasibility, not semantic use, latent scene induction, or explicit scene construction. The next tests isolate history use, compare aligned with permuted targets, and introduce varied lexical prototypes before video and multiscale composition; coherent responses to scene edits remain a later criterion.

Keywords Language Modeling, Scene-Mediated Generation, Latent Scene Induction, Multiscale Composition, Persistent World State, Concept Bottlenecks, Grounded Cognition, Mental Imagery

1 Motivation

Autoregressive language models already combine the preceding sequence into distributed contextual states. Controlled studies show that predictive training can produce causally active representations of entities, bindings, board states, space, and time; in one controlled study, however, entity tracking did not arise reliably from text pretraining alone and degraded on harder splits [1, 2, 3, 4, 5]. STLM therefore does not begin from the claim that ordinary models lack internal world structure. It asks whether an externally organized intermediate, made compulsory for each lexical decision, can systematically shape and expose that structure.

In STLM, the substrate is constitutive rather than optional: every lexical decision must pass through it. A *substrate predictor* (SP) emits an externally structured image, description, video, or related representation. A *token readout* (TR) must recover the next word from substrate states alone, without a token-side context channel. The intended gain is a

different inductive bias: properties such as color, object identity, motion, containment, contrast, and causation recur across many targets. The central question is whether this shared organization changes the computation that supports prediction rather than merely replacing token IDs with a more expensive codebook.

1.1 From Contextual Integration to Scene Emergence

We use *scene* inclusively for a joint configurational state containing physical entities and spatial relations, but also events, roles, mechanisms, constraints, and abstract schemas. A couch contributes an object with spatial requirements; "however" contributes a contrast relation; a promise contributes agents, obligations, and possible violations. Pixels and video are natural carriers for physical structure, while structured descriptions can preserve mechanisms, relations, and consequences where literal depiction is strained. The intended representation is broader than visual imagery but more constrained than an arbitrary hidden state: its authored organization must support testable composition and intervention.

The motivating analogy is human reading, in which prose can prompt a reader to construct and revise a scene rather than retain only surface wording. Perceptual-symbol, dual-coding, conceptual-metaphor, and embodied accounts provide precedents for relating linguistic, imagistic, and configurational representations [6, 7, 8]. STLM neither assumes that all reading is visual nor attempts to reproduce human imagery, and its architecture does not depend on those cognitive accounts being correct. They motivate the choice of substrate; causal controls must establish what the model actually uses.

Entangled Alignment names a related corpus-level omission *the Missing Reader*: source text records conclusions more fully than the reader-side questions, expectations, evaluations, and revisions through which understanding changes [9]. STLM isolates a complementary candidate omission. Text corpora generally do not preserve the visual or schematic state that a reader might construct while following a story, mechanism, proof, or program. The proposed substrates are authored synthetic representational targets, not recovered human mental imagery; whether any such target helps a model is an empirical question.

Three levels must be distinguished. *Ordinary contextual integration* is the combination of prefix information that any autoregressive transformer may perform. *Latent scene induction* is the hypothesis that repeatedly predicting aligned, structured targets changes this integration so that reusable physical or schematic factors coexist and interact, even before one explicit persistent scene exists. In Stage 1 this could occur inside the SP or in a causal readout over its emitted trace; previous substrate states are not fed back into the predictor or explicitly assembled into a scene. *Explicit scene construction*, proposed for Stage 5, instead maintains one revisable substrate state across the discourse.

Shared meaning could matter through factor reuse. Fifty thousand unrelated codewords require fifty thousand arbitrary associations. Aligned targets may instead reuse redness across objects, containment across scenes, paths across motions, opposition across arguments, and agent–action–consequence structure across mechanisms. A capacity-limited predictor may find such factors cheaper than memorizing every association independently. Fixed canonicals still admit a codebook shortcut, so this pressure is a hypothesis rather than an architectural consequence.

For visual substrates, a medium-respecting generative decoder could add *rendering pressure*: its emitted representation would have to support reconstruction of boundaries, depth, containment, material, motion, and occlusion. This may favor reusable visual factors but cannot guarantee them. The present dense fixed-canonical pilot does not impose this pressure (§8.4).

Multiscale targets would add consistency across representational scopes. A stable ball primitive, a green attribute, and a rolling trajectory could recur in a phrase or sentence target depicting the same green ball rolling; a paragraph-level target could preserve that object's identity as the discourse changes it. Cross-scale targets would therefore be constructed from a shared registry rather than authored independently. This is intended to make factor reuse advantageous without guaranteeing that the model will choose it over a private code.

Informally, the proposal resembles an *intuitive chain of thought* that is non-lexical and, in visual variants, spatial. Unlike verbal chain-of-thought, it adds no optional rationale or deliberation tokens: the substrate lies on the compulsory route to each word. The SP's internal coordinates need not be intelligible to a human, but its emitted substrate must remain decodable within the authored representational space.

The operative tests are causal rather than merely visual. Disrupting intended structure should selectively impair relevant predictions, while controlled semantic edits should produce corresponding changes. Scene emergence requires three capabilities under these tests: *composition*, combining familiar factors in novel configurations; *manipulation*, changing a relation and propagating its consequences; and *reuse*, applying the same operation across unrelated content. A private code that merely mediates prediction does not qualify. This computational definition makes no claim about consciousness or phenomenal imagery.

In the persistent version, each candidate update must remain compatible with the accumulated state. An object occupies space and preserves identity; an event constrains later consequences; a contrast relation requires two propositions in tension. Training against continued language would penalize incompatible updates when they make later words harder to predict. This *coexistence filter* is the proposed mechanism for coherence: later predictions can use structures maintained together rather than reconstructing every relevant relation from lexical associations. Stage 5 tests whether making such a state persistent, externally shaped, and obligatory changes prediction.

1.2 Research Trajectory

The program varies two partly independent dimensions. *Compositional scope* runs from lexical items through phrases, sentences, paragraphs, and evolving discourse. *Temporal medium* runs from still configurations through short transformations to continuing dynamic state. A sentence-level still scene can test composition before word-level video is useful, while video can improve the representation of an event without solving discourse-level persistence.

Along the scope axis, the weakest intervention assigns one fixed canonical to each lexical item. Context-selected prototypes situate that item within the prefix; multiscale targets preserve components across lexical and longer-span representations; explicit persistence maintains entities, attributes, relations, events, and schemas across discourse. Along the medium axis, still images and structured descriptions represent configurations, while video represents transformations and trajectories. Scope and medium can therefore be varied separately.

The program is easiest to understand as five developmental stages. *Stage 1*, implemented here, assigns one fixed still image or encoded description to each word. *Stage 2* gives a word many valid instances, first by random sampling and then by prefix-conditioned selection, so that the emitted substrate can represent the word as used in that context. *Stage 3* replaces the still lexical instance with a short clip when motion or transformation matters. *Stage 4* lets one image, clip, description, or schematic state represent a phrase, sentence, paragraph, or multiscale bundle. *Stage 5* maintains one scene continuously across discourse and requires lexical output to be read from that evolving state.

This ordering is a communicative roadmap, not a strict dependency graph. Scope and medium remain partly independent: a sentence-level still can be tested before word-level video, and video does not by itself create persistence. Later machinery is warranted only when the simpler intervention fails to produce the predicted causal effects. Stage 4 makes exact span decoding a distinct open problem because one scene can support many word sequences.

At Stage 1, the single-state readout TR-S uses only the present emitted state, while the causal readout TR-C uses the complete emitted 128-state window through the present position. Neither is the persistent scene of Stage 5: TR-C reads a trace, but it does not revise one accumulated world. Section 2 specifies this compulsory seam in detail.

1.3 Empirical Questions and Evidence Boundary

The empirical questions follow the two axes. Can a compulsory substrate carry next-token information, and do its emissions preserve intended structure rather than only machine-readable identity? Do aligned targets outperform fixed permutations and random targets? Do multiscale targets improve held-out composition and binding beyond matched word-level targets; do clips improve temporally structured prediction beyond alignment-validated stills; and does explicit persistence improve coherence and counterfactual manipulation beyond nonpersistent substrate sequences? Gains should be structure-specific rather than uniform across benchmarks; §7 states the predicted profile and its falsifying signature.

The pilot implements two media (§8). The *pixel substrate* is a literal 64×64 RGB image with an explicit (x, y) axis. The *text-description substrate* is the frozen text encoder's full hidden-state sequence for the same authored description. Both carry decodable predictive signal, with descriptions outperforming pixels under fixed-canonical regression, but this comparison ranks two complete pipelines rather than isolating geometry or meaning.

The paper separates four claims. The *operational claim* is that a substrate-only readout can recover context-dependent next-token information; the five-configuration pilot tests this. The *alignment claim* is that intended semantic organization matters more than an arbitrary assignment, which requires matched aligned, permuted, and random targets. The *fidelity claim* is that emitted states remain on the authored configurational manifold rather than carrying token identity in artifacts. The present pilot tests neither alignment nor fidelity conclusively.

The *scene-emergence claim* is stronger. Alignment-dependent compositional transfer alone does not establish a scene-like state. Latent scene induction additionally requires evidence that multiple bound constituents coexist, retain or update across positions, and respond appropriately to causal edits or patches. Explicit scene construction further requires direct intervention on one persistent temporal state. Neither form is demonstrated here.

The individual ingredients have precedents in concept bottlenecks, continuous-representation language models, visual augmentation, latent reasoning, and persistent world-state systems (§6). The proposed conjunction is an externally organized substrate made compulsory at every lexical decision, matched interventions on its semantic alignment, cross-scale reuse of stable components, and eventually one persistent physical-and-schematic scene. The present pilot demonstrates only the prerequisite interface.

Contributions. The paper (i) defines STLM's compulsory lexical seam; (ii) separates ordinary contextual integration, latent scene induction, and explicit scene construction; (iii) gives a five-stage development path across compositional scope and temporal medium and specifies causal tests for each claim; and (iv) reports a reproducible, two-medium, five-configuration stage-1 pilot over a 1024-token TinyStories vocabulary.

2 The Lexical and Multiscale STLM Architecture

Stages 1–4 share two networks, both frozen at inference, separated by an explicit *intermediate substrate*. The toy in §8 is the restricted Stage-1 fixed-canonical regression specialization. The *substrate predictor* (SP; large, the imagery engine in visual media) reads running text and emits a configurational representation. The *token readout* (TR; small, the pen) maps substrate states to tokens, either from the current state alone (TR-S), from the causal emitted sequence through that state (TR-C), or in Stage 4 from one substrate object to a variable-length span. No readout receives source-token embeddings, an SP hidden state, or a parallel text pathway. A third artifact, the vocabulary V, is constructed once (§3) and shapes both networks; pixels, encoded descriptions, or other configurational representations can instantiate its medium.

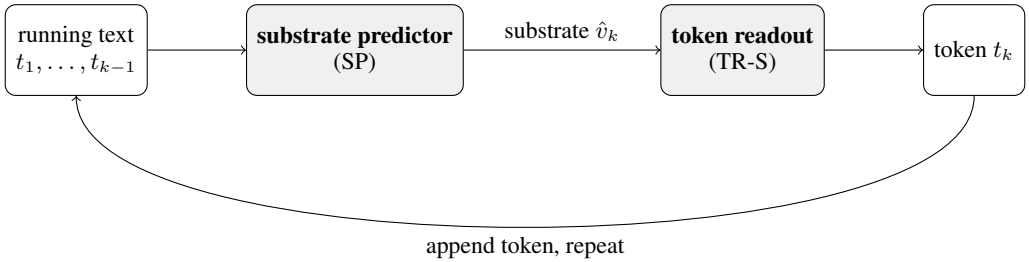


Figure 1: The stage-1 TR-S loop. At each position the substrate predictor (SP) reads the running text and emits a substrate state; TR-S names it as one token; the token is appended and the loop continues. The only channel between SP and TR-S is the current substrate state. The TR-C alternative freezes the same SP and trains a fresh causal readout over the complete 128-state window through the current position, still without source-token access.

The substrate vocabulary V. A frozen mapping from token IDs to sets of vectors in $\mathbb{R}^d$. For each token t, $V(t) = \{p_1^t, \ldots, p_n^t\}$ is a set of *semantic prototypes*: images, video clips, encoded descriptions, or reconstruction latents that depict or encode plausible instances of t's referent. For concrete tokens these may be many images or clips of the object or event; for abstract or function tokens they are LLM-authored image-schema or metaphorical instances. We write $\bar{V}(t)$ for the centroid, a summary statistic used only for analysis rather than as the training target (§3). The choice of encoder E is the program's most consequential pre-training decision: the substrate inherits whatever geometry E chooses to preserve. §9 returns to the choice between contrastive and reconstruction-objective encoder families.

A token's prototype set can be used in three increasingly strong ways. Stage 1 stores one canonical per word, making the SP a plain regressor to a fixed target. Stage 2 first uses a *random semantic-prototype* policy that samples a fresh member of $V(t_k)$ whenever the next token is t_k, forcing the SP and TR to learn the whole semantic cloud rather than one point. A stronger Stage-2 *context-selected semantic-prototype* policy chooses, for each training instance, the member of $V(t_k)$ that best matches the prefix context: the green apple in an orchard context, the bitten apple in a lunch context, the corporate-logo instance in a technology context. The selection step may use the prefix and the supervised target token during training, but it must not use future text or any side-channel unavailable to the SP at inference. Random selection tests whether the architecture can learn a word's prototype distribution; context-selection tests whether the substrate can carry the situated meaning of this particular occurrence. These are different experiments and should not be conflated.

Stage-4 multiscale substrate composition. The per-token vocabulary is the smallest member of a possible hierarchy. Longer-span targets can depict or encode phrases, sentences, and paragraph-level states while reusing an explicit registry of the entities, attributes, motions, relations, and abstract schemas used in their lexical components. One

compulsory-seam implementation predicts a bundle at every position: a next-token substrate, one or more longer-horizon scene hypotheses, and a discourse-level scene hypothesis, all from the same text prefix. The TR receives only this predicted bundle; the bundle is recomputed rather than carried forward and revised after the token is appended. During training, longer-span targets may be constructed from the continuation being predicted, but they are targets rather than inputs and are unavailable at inference except through the SP's prediction. A variable-length span decoder is the more direct formulation when one scene is intended to represent one complete semantic span (§9). Neither extension is implemented in the pilot.

Substrate predictor (SP). A large text-conditioned generative model. In the single-scale design, its input is the text prefix with full long context and its output at every position is one substrate state $\hat{v}_k \in \mathbb{R}^d$ *sampled* from its learned distribution over plausible next-word representations. Architecturally it is a causal transformer with a generative head — diffusion noise prediction in $\mathbb{R}^d$, or autoregressive prediction over a discretization of $\mathbb{R}^d$ — not a regressor at frontier scale. §3 explains why this matters: a regressor trained against a multi-prototype set collapses to its conditional mean, washing out the per-instance structure the substrate is meant to carry; a generator trained on sampled or selected prototypes learns the distribution and emits crisp instances. After training the SP is frozen.

Token readout (TR; the pen). A learned function that maps the SP's emitted substrate representation into a distribution over tokens. TR-S is a stateless map $f_{\mathrm{S}} : \mathbb{R}^d \to \Delta^{|\mathrm{vocab}|}$ invoked once per generated token; it receives only the current per-position substrate $\hat{v}_k$. In the reported toy comparison, TR-C is a causal sequence map: a medium-specific state adapter, shared over time, independently maps each $\hat{v}_j$ to 256 dimensions; a two-layer, eight-head causal Transformer maps $(\hat{v}_1, \ldots, \hat{v}_k)$ to vocabulary distributions at every aligned position. Its position-k decision can depend on substrate states $\hat{v}_{\leq k}$, including the current state, but it receives neither the original token prefix nor an SP hidden state. TR-C is freshly initialized and trained as a complete readout rather than as an adapter to TR-S. For one grid state TR-S can be a CNN; for a description state it can be a small transformer over the substrate's internal positions.

In the current Stage-1 design, the SP integrates the text prefix through an ordinary causal Transformer but does not consume its earlier emitted substrates. TR-C is the only component that directly receives their sequence. It can therefore integrate information across multiple substrate states, but its trace is not an incrementally revised scene and the architecture alone does not imply that it combines them according to their intended meaning.

Conceptually, the division of labour mirrors §1: the SP has *already produced* the substrate representation, and the TR performs the modality conversion into the lexical form the output stream requires. A TR is functionally specialized to its own SP's substrate idiom rather than being a generic decoder of arbitrary states; §3 separates its training regimes. The architectural constraint is that no TR receives a text-context side channel. For TR-S, predictive sufficiency must rest on the current SP output; for TR-C, it may rest on the causal substrate sequence through the current state.

The readout seam. In Stages 1–3 and the Stage-4 bundle formulation, one substrate object and one readout invocation occur per generated token. In the direct Stage-4 formulation, one object and one invocation instead produce a variable-length span. No emitted object is revised before Stage 5. This constrains the interface, not the SP's internals: a diffusion or autoregressive head may run many iterations before emitting it. Each $\hat{v}_k$ is understood as a *positionally or configurationally structured* representation — a pixel grid, a video clip, an encoded-description sequence, or another object with internal positional and relational content — rather than one global compression. Spatial-medium claims additionally require an explicit spatial axis; §9 elaborates.

3 Training Lexical STLMs: Stages 1–2

Stages 1–2 separate construction of the intended substrate organization from learning the two sides of the lexical seam. Phase zero builds and freezes the substrate vocabulary; phase one teaches the token readout to recognize states from that vocabulary; and phase two teaches the substrate predictor to generate an appropriate next-token state from linguistic context. Stage 1 uses one fixed canonical per word; Stage 2 uses variation within each word's semantic prototype set. Keeping these roles distinct makes it possible to study readout fidelity, adaptation to predictor outputs, and optional end-to-end coupling as separate design choices.

3.1 Phase Zero: The Substrate Vocabulary

The substrate vocabulary is computed once before model training. Stage 1 authors one description and derives one canonical target per token. Stage 2 instead generates many varied descriptions (for concrete tokens, scene contexts such as "a red apple on a wooden table"; for abstract tokens, metaphorical descriptions via an LLM-as-dictionary step; for function tokens, glyph or substitute-word descriptions). Phase zero uses three distinct maps. A description substrate

applies a frozen *text* encoder, $d_t \mapsto E_{\text{text}}(d_t)$. A pixel substrate applies a text-conditioned *generator*, $d_t \mapsto G(d_t) = I_t$. A latent substrate additionally applies an *image* encoder to the rendered result, $I_t \mapsto E_{\text{image}}(I_t)$. A text encoder produces a conditioning representation rather than an image or spatial image latent. Cache the resulting canonical or prototype set and, for Stage 2 analysis, its centroid per token.

At frontier scale, a 50k-token vocabulary with 500 prototypes per token would contain 25 million prototypes, making rendering, validation, storage, and deduplication substantial design costs. The resulting vocabulary file is read throughout training and retained only as a reference afterward; neither the phase-zero auxiliary LLM nor the encoder is retained after construction. Section 8 reports the measured construction cost of the pilot.

3.2 Phase One: Training the Token Readout

Standard multi-class classification on the vocabulary alone, with no text corpus. Train the TR with cross-entropy to output the correct token from a (substrate, token) pair. In Stage 1 it sees the single canonical for each token; in Stage 2 it samples across the prototype cloud and learns within-word invariance. In phase two (§3.3) it is then either frozen, continued in lockstep on freshly sampled prototypes, or trained directly on the SP's actual outputs — three design points along the readout-training axis, trading off convergence stability against the human recognizability of the SP's emitted substrates. The TR-S/TR-C distinction is orthogonal to this axis: either scope can be trained on actual SP emissions, but TR-S learns from individual states whereas TR-C learns from causally masked sequences with an aligned loss at every position.

The choice also decides what the TR can exploit, and three mechanisms must be kept apart. When a TR is trained after a frozen SP on that predictor's actual emissions — the option the toy of §8 takes — it can adapt to stable deviations and artifacts in the SP's output distribution. Because the SP is frozen and no gradient crosses the seam, this is *one-way decoder adaptation*, not mutual co-adaptation: the readout moves toward the predictor, never the reverse. *Joint private coding*, in which the pair settles on a shared convention that need not respect the targets' meaning, becomes possible only under an end-to-end objective in which readout gradients reach the SP — the optional coupling term §3.3 reserves. A third regime trains the TR only on authored canonicals, making it a *fidelity diagnostic*; it constrains the SP only if its loss is included in the SP's objective. Under the no-gradient scheme specified here it is a measuring instrument rather than a training pressure. This independence creates distribution shift because a canonical-only reader never sees SP drift. Meaningful within-concept variation (§9), rather than exact canonical points, is therefore required for a semantic reader.

The TR is intended to be small (1–10M parameters) and is never exposed to source tokens directly; at toy scale TR-S trains in minutes, though a 50,000-way classifier over 25 million prototypes would not. The restriction still permits both readouts to learn token priors and exploit label-correlated artifacts because they see token labels and their empirical frequencies during training. TR-C can additionally model regularities across a causal substrate sequence, including language regularities encoded there indirectly. The architectural guarantee is the absence of a token-side context pathway, not the absence of all lexical statistics or contextual information.

Why a learned token readout rather than nearest-prototype lookup. SP emissions can deviate systematically from stored prototypes, so nearest-neighbor decoding is restricted to the chosen metric and the canonicals' Voronoi boundaries. A learned TR can absorb those deviations and learn invariances from actual SP emissions. That flexibility also creates an artifact channel: a TR trained on SP outputs may exploit stable cues unrelated to authored meaning. Section 9 therefore proposes comparing an adaptive channel readout with a frozen manifold readout.

3.3 Phase Two: Next-Substrate Prediction

Conditional training over the substrate vocabulary. For each position k, the policy defined in §2 supplies a target $p_k \in V(t_k)$. The substrate predictor reads the prefix $t_1, \ldots, t_{k-1}$ and is trained against p_k. Stage 1 uses the fixed canonical with the pilot's regression or cosine objective. Stage 2 uses a generative objective, such as diffusion noise prediction or autoregressive cross-entropy over a discretization: random selection teaches the full conditional distribution $p(\hat{v} \mid t_{1:k-1})$ over a word's prototype cloud, while context selection teaches the contextually appropriate instance.

This objective choice is consequential. MSE regression onto a fixed target (the centroid), or even onto resampled prototypes, collapses to the conditional mean — a probability-weighted blur of all plausible next-word images. A generative objective with resampled targets matches the distribution rather than its average, so the predictor's emitted substrate is a crisp specific instance rather than a smeared superposition. This matters most for relational and abstract terms, whose per-instance visual variation is largest and whose means are therefore the most washed-out.

What fixed-target regression computes. The collapse has an exact characterization, and stating it also bounds what a feasibility result can mean. Let $c_y \in \mathbb{R}^d$ be the canonical target for token y, and let C be the matrix whose rows are

those canonicals. A predictor minimizing squared error against c_Y has Bayes-optimal deterministic output

$$g^*(x) \;=\; \mathbb{E}[c_Y \mid X = x] \;=\; \sum_y p(y \mid x)\, c_y \;=\; C^\top p(\cdot \mid x),$$

the posterior-weighted average of every plausible next word's canonical. Four consequences follow. First, blurring is not an implementation defect but the objective's optimum: under next-token uncertainty the best fixed-target output *is* a superposition. Second, if C has full row rank the posterior is in principle recoverable from $g^*(x)$ by a linear map, so a readout can succeed even when the canonicals carry no meaning at all. Third, seam feasibility is therefore weak evidence by itself: the pixel substrate offers 12,288 coordinates and the flattened description substrate 21,504 for only 1024 classes, so either target set has ample room to act as an error-correcting output code [11]. Fourth, semantic alignment cannot change what is expressible through such a channel; it can only change the inductive bias — which functions are cheap to learn, how they generalize, and how they transfer compositionally. This is precisely why the central claim is posed as a question about sample efficiency and transfer under a permutation control, rather than about whether the seam can carry information.

This exact linear characterization applies to the pixel condition, whose loss is squared error on unconstrained outputs. The description condition instead applies a per-position cosine loss to unit-normalized outputs, so its optimum at slot j is the *normalized direction* of the corresponding conditional mean, $g_j^*(x) \propto \sum_y p(y \mid x)\, \bar{c}_{y,j}$ with $\|g_j^*(x)\| = 1$. Because each slot discards magnitude separately, full row rank of the raw flattened description matrix no longer guarantees that the token posterior is linearly recoverable, even though the mapping can still carry substantial class information.

Under the optional lockstep regime, the TR is trained on (p_k, t_k) in the same step and against the same sampled target used by the SP, with one optimizer per network and no gradient crossing the seam. The shared prototype is a data anchor rather than an optimization coupling: the SP sees p_k as a generative target, while the TR sees it as classification input. Lockstep training also permits a later end-to-end coupling term, $\mathrm{CE}(f(\hat{v}_k), t_k)$, without restructuring the pipeline. We call the objective *next-substrate prediction*; Stages 1–2 specialize it to still images or descriptions and Stage 3 to video. At deployment, the frozen SP and TR run in series; no training-only component participates in inference.

4 Proposed Stages 3–4: Temporal and Multiscale Substrates

4.1 Stage 3: Lexical Video Substrates

Stage 3 enriches the substrate from still images to short videos. The architecture's structural shape is unchanged — the SP reads text, emits a substrate state, and the TR names it as a token — but the substrate at each position is now a short video clip rather than a single image.

Many words are not static. "Threw" is a trajectory — a grasp, a wind-up, a release, a flight, an arrival; "melted" is a solid-to-liquid transition; "collapsed" is a process with a before and an after. Even "door" carries the dynamic of opening and closing, while "bridge" carries the affordance of crossing. A still image can capture only one moment of this time-structure; flattening a verb, process, or transition into a single configuration is intrinsically lossy in a way depicting a roughly time-invariant table is not. A video substrate instead represents trajectories and transitions across frames — the throw's motion, the melting, the door's swing — so its expressive capacity better matches what these words encode.

The vocabulary stores short video clips per token or, more practically, sequences of vectors in a video encoder's geometry. Vocabulary construction uses a text-conditioned video *generator* to render each clip and that generator's autoencoder to represent it — the components of systems such as Stable Video Diffusion [12] or Sora [13] — rather than an image pipeline.

The empirical question is whether extending the substrate's temporal capacity from one moment to a clip-sized window measurably reduces the failure modes that still images cannot handle. Stated precisely, Stage 3 tests whether *temporally structured targets improve temporally structured prediction*. Stages 1–3 supervise each position against an independently chosen next-word target, so none enforces substrate-level coexistence. Aligned targets might nevertheless induce a latent scene-like precursor inside contextual computation; that possibility must be demonstrated through matched semantic controls rather than inferred from architecture. The empirical prediction is that Stage 3 outperforms matched still-image Stage 2 on tasks where verb aspect, action sequencing, and short-range temporal trajectory matter, while remaining comparable on tasks dominated by static configurations. Because temporal medium and compositional scope are independent, the same comparison can later be repeated with phrase- or sentence-level targets.

4.2 Stage 4: Span-Level Substrates

Stage 4 increases semantic scope rather than requiring a new medium. One image, video, description, or schematic state represents a phrase, sentence, paragraph, or other meaning-bearing span while reusing the same entities, attributes, relations, and motions as its lexical components. These span objects remain independently predicted; unlike Stage 5, they are not revisions of one state that persists through the whole discourse.

The direct formulation predicts one substrate object and decodes a variable-length source span from it. A multiscale alternative predicts lexical, phrase-level, and longer-horizon objects together at each position and reads the next token from that bundle. The direct formulation tests whether one joint scene can support a coherent stretch of language; the bundle formulation preserves exact per-token timing while testing whether cross-scale factor reuse helps prediction. Both require matched controls because a high-dimensional span object can still become an arbitrary code.

Exact wording is not uniquely determined by a scene. Stage 4 can therefore incur high token loss even when a decoder produces a valid description of the represented event. This is the same one-observation-among-many problem faced by ordinary language modeling, but it is stronger after linguistic detail has been compressed into a scene. Evaluation should retain likelihood while also measuring semantic adequacy, consistency with controlled scene edits, and whether generated spans preserve entities, relations, and consequences.

5 Proposed Stage 5: Persistent Scene Construction

Stage 5 externalizes and strengthens the scene-induction hypothesis. Instead of relying on independently predicted per-position or per-span targets, free-running generation maintains one continuously updated configurational scene-state across the discourse, and the token readout receives only a state selected from it. The scene is neither an optional memory retrieved when useful nor a sequence of independently predicted pictures. Each generated word updates the same working state; later lexical decisions must therefore be compatible with what has accumulated. This is the proposed *coexistence filter*. Stage 5 has not been implemented, and the stage-1 pilot contains neither scene accumulation nor scene-level compatibility pressure.

Operationally, Stage 5 retains STLM's two serial networks. At lexical position k, the SP transforms the accumulated scene S_{k-1}, conditioned on the generated prefix, into an updated state S_k that may contain one or more newly predicted frames. The TR receives only S_k and emits the exact next-token distribution. That token is appended, the scene is updated again, and the loop repeats. A single shared network that instead receives completed text-and-scene history and predicts the next scene interval and next text together is a simpler joint world–language architecture, not the compulsory two-network STLM.

Training can pair each source event or span with the corresponding interval of an authored or generated scene trajectory, then derive successive scene states for lexical supervision. The signal is underdetermined: one scene can license many accurate word sequences, and literal video may omit syntax, emphasis, abstract relations, and discourse intent. A Stage-5 substrate may therefore combine visual change with schematic structure rather than consist of naturalistic pixels alone.

The substrate medium and form should remain experimental variables throughout the program. A human-intuitive or photorealistic scene need not be the representation most useful to a model. Literal scenes, moving diagrams, equation states, execution traces, symbolic-spatial layouts, and hybrid states should be compared on matched tasks rather than selected by visual appeal. The companion proposal *Imagining the Corpus* treats this as an upstream problem: compare candidate forms under matched task, capacity, compute, and supervision; retain those that improve over the matched baseline and survive alignment and intervention controls; and only then scale persistent trajectories across documents [10].

High exact-token loss would count against Stage 5 as a language model even when its generated continuation is fluent and scene-consistent. Because one recorded continuation does not exhaust the valid descriptions of a scene, likelihood should be reported alongside semantic faithfulness, persistent-world consistency, controlled counterfactual response, and grounded inference-time utility. These complementary measures do not replace the language-model criterion.

The intended scene carries both *physical/visual content* (the running boy, the kitchen layout, the light) and *schematic/configurational content* (relations, abstractions, image schemas [8] — containment, support, path, balance, force, contrast) in one working state. The distinction is functional rather than a literal architectural partition. Both kinds of content are subject to the same compatibility pressure: a candidate "however" must fit the contrast structure already present, while a candidate "couch" must fit the spatial structure.

Scene contents and organization. The proposed working state represents objects, layout, motion, and light, and updates them as the discourse unfolds: a boy crosses the room, the light shifts, a door opens. It need not render pixels at inference, but its physical factors must remain anchored to a decodable spatial representation. Coexistence supplies the proposed constraint: a placed object occupies space, occludes what lies behind it, and retains material properties, so incompatible later updates should make continued language harder to predict.

The state would also need to represent relations and abstractions without straightforward visual instantiations. A contrastive term such as "however" could correspond to a relation between propositions; "the argument collapsed" could reuse a transformation schema without requiring literal imagery. Flow, path, and balance are candidate structures for concepts such as inflation, time, and equality. Whether such reuse emerges, composes with physical content, or causally affects prediction is an empirical question.

At each position, a learned readout would have to select the information relevant to the next token without access to an unrestricted linguistic bypass. Concrete and abstract terms need not occupy predefined regions, but interventions should reveal whether distinct physical and schematic factors are represented and used. How the readout performs this selection remains an open Stage-5 design problem.

Physical and schematic content may be distributed, superposed, nonlinearly separable, or related by an arbitrary rotation within one state. We hypothesize that a joint objective anchoring physical structure to reconstruction and schematic structure to relational language would give the two functionally distinguishable roles. Probe separability alone would not establish such specialization; the tests under H5 (§7) therefore require causal intervention.

What Stage 5 tests. Stage 5 predicts gains where generation depends on maintaining and manipulating a shared state: entity and spatial consistency over long spans (a character introduced earlier remains present unless an event moves or removes them), continuation after controlled counterfactual edits (remove an acquired key and propagate the consequences), and continuity of abstract relations. The gains should be selective, survive matched-compute comparison, and follow the direction of interventions on the persistent state. Section 7 states the corresponding tests.

Open design problems. Implementing Stage 5 requires machinery for maintaining and updating the scene across positions, analogous to persistent video world models. A concrete architecture must specify how event-level scene intervals synchronize with variable-length language; how the readout selects relevant state without an unrestricted hidden-state bypass; which loss prevents drift into an arbitrary recurrent code; whether physical and schematic factors acquire causally distinguishable roles despite being distributed or superposed; how stochastic prediction defines token probabilities; and whether iterative substrate generation can remain synchronized with lexical or span emission at acceptable inference cost. These questions define the central technical agenda for Stages 3–5.

6 Related Work

The relevant prior work is best separated along three axes: whether a substrate representation is an *obligatory route* to lexical output, whether generation maintains an *evolving scene or world state*, and whether the state's *semantic organization* is tested causally. Prior systems realize important subsets of this program. Stage 1 isolates the compulsory per-token seam; the proposed Stage 5 combines that seam with one persistent physical-and-schematic scene and matched semantic interventions. To our knowledge, no prior free-running natural-language system combines all three.

6.1 Concept Bottlenecks and Representation-Space Language Models

Concept Bottleneck Models first predict human-specified concepts and then make the final task prediction only from those concepts [14]. This is the clearest architectural precedent for an obligatory, inspectable seam. CB-LLM applies the pattern to large text-classification benchmarks using automatically generated textual concepts [15]; CB-pLM places an interpretable concept layer before token prediction in a generative masked protein language model [16]. These systems establish both the two-stage commitment and its application to language-model outputs, although they do not study left-to-right natural-language generation through per-token visual or description canonicals. Subsequent work showed that CBM intermediates leak task information beyond the intended concepts [17]; the channel-versus-meaning distinction that motivates our controls is the generative analogue of that finding.

Large Concept Models (LCMs) are the closest broad precedent [18]. An LCM autoregressively predicts the next sentence in a fixed SONAR embedding space and uses a frozen SONAR decoder to reconstruct text. It therefore already demonstrates language generation through an obligatory meaningful representation and a decoder that does not receive the original token context. Continuous Autoregressive Language Models similarly predict reconstructable continuous vectors for multi-token chunks [19]. At token granularity, Projected Autoregression makes continuous next-token

vectors the primary autoregressive interface, studies how embedding geometry affects stability, and commits outputs by nearest-neighbor projection [20]. Earlier continuous-output sequence models already predicted word embeddings in place of a softmax [21] or used error-correcting output codes for sequence prediction [11]; these are important alternative explanations for why structured continuous targets can help, independent of any configurational reading. Next-concept prediction learns discrete multi-token semantic units and predicts them alongside ordinary tokens [22], sharing the goal of routing prediction through semantic units while retaining the token pathway.

Continuous and concept-space language generation are established precedents: LCM supplies obligatory sentence-space generation, while Projected Autoregression studies learned or pretrained token-embedding geometry and commitment dynamics. STLM's narrower proposal uses externally authored modality-derived organization, extends it toward a persistent scene-state, and holds the target set fixed while permuting word-to-substrate alignment to test meaning causally.

6.2 Visual Supervision and Machine Imagination for Language

Vokenization contextually associates each language token with a related retrieved image and uses the resulting "vokens" as visual supervision [23]. VaLM retrieves images for each textual context and fuses their representations into an autoregressive language model before the ordinary vocabulary projection [24]. These are strong precedents for token-level visual grounding, but the text hidden state remains available to the output head.

Machine-imagination systems move still closer to the motivating intuition. iNLG synthesizes an image from the input context and supplies it as a visual prefix for open-ended text generation [25]; LIVE synthesizes images at sentence granularity and fuses selected image features into BART or T5 [26]. These systems use images for a related reason: visual structure may supply information and organization that text alone does not make easy. In both cases, however, the image augments rather than replaces the textual route to generation, so the language model can ignore it. BLIND-VALM reports that visually grounded text representations can match explicit image retrieval on its evaluations [27], an especially relevant precedent for comparing STLM's pixel and description media. STLM's no-bypass design makes the medium's causal contribution measurable.

6.3 Scene Construction and Persistent State for Language

TALE-SPIN simulated physical and social worlds before rendering their events as stories [30]; modern narrative systems maintain dynamic entity states that guide subsequent generation [31]. Imagine-and-Verbalize constructs a relational scene knowledge graph before verbalization and, for stories, alternates graph and sentence generation [32]. These systems establish that explicit evolving or scene-structured state can guide language. Their state is symbolic or entity-centered, however, and their verbalizers retain ordinary linguistic routes or operate at plan, sentence, or task granularity rather than making one learned physical-and-schematic scene the sole interface at every token.

Closer to persistent imagery, Machine Mental Imagery incrementally creates and edits external visual artifacts during situated dialogue and later retrieves that history for responses [33]. Orca learns a multimodal world latent with prediction and modality-specific readouts, while language answers still use the model's ordinary VQA language head [34]. These are central precedents for persistent scenes and world-state readout. They do not, however, recurrently build one scene from free-running generated text and require every lexical output to be read only through it. That exact conjunction is the unimplemented Stage 5 proposal.

6.4 Latent Situation Models and Continuous Thought

Structured state can arise without explicit scene supervision. Li et al. identify contextual representations in BART and T5 that track entities' changing properties and relations, and show that interventions on those representations predictably alter generation [1]. In a synthetic setting, a GPT trained only on Othello move sequences develops a board-state representation whose interventions control move predictions [2]. Natural-language models also exhibit transferable mechanisms for binding attributes to contextual entities [4]. Linear probes reveal multiscale representations of geographic space and historical time, while neuron ablations and a targeted time-neuron intervention provide narrower evidence that some identified features affect prediction [5]. These results make latent scene induction plausible, but they do not establish a unified dynamic scene in ordinary language generation or show that an externally authored substrate causes the organization.

Coconut provides a direct precedent for continuous computation without decoded intermediate tokens: during a designated latent interval it feeds the language model's last hidden state back as the next input embedding [35]. Its continuous thoughts are learned end-to-end for downstream reasoning rather than anchored to an externally specified visual or schematic organization, and lexical generation resumes through the ordinary language-model head. STLM

instead constrains the representation used at ordinary lexical positions. Its proposed intuitive chain is compulsory at those positions rather than a designated iterative rollout.

6.5 Text-Conditioned Image and Video Generators and Their Latent Spaces

Text-conditioned generators combine a conditioning representation with a generative model operating over pixels or spatially structured visual latents, and the two components should not be conflated. Imagen [36] shows that a *generic text-only* T5 encoder provides effective conditioning for image synthesis — which does not make the T5 representation itself a spatial image representation, since the visual structure is learned by the diffusion model rather than carried by the conditioning embedding. Latent Diffusion [37] and FLUX [38] operate over compressed spatial image latents; Stable Video Diffusion [12] is a latent *video* diffusion model, and Sora [13] generates over compressed spacetime latents with a corresponding decoder. For this program the candidate substrates are therefore pixels or the spatial latents of an image or video autoencoder; text-conditioning embeddings are a distinct substrate family without an explicit spatial grid. Our vocabulary construction applies these systems to a downstream language-modeling task rather than to visual generation.

6.6 Joint-Embedding Predictive Architectures

Joint-Embedding Predictive Architectures (JEPA) [39] predict in representation space rather than reconstructing raw observations. I-JEPA [40] and V-JEPA [41] learn visual or video representations by predicting the embeddings of masked or future parts of an observation; V-JEPA 2 extends the family toward prediction and planning in embodied settings [42]. Language-adjacent variants such as LLM-JEPA [43], VL-JEPA [44], and JEPA-T [45] show that the broad JEPA commitment — predict useful target representations rather than raw outputs — is already entering text and vision-language systems.

STLM adopts JEPA's broad move of predicting a target representation rather than directly predicting a token. JEPA systems primarily learn useful representations for downstream tasks, while STLM is an autoregressive language model whose deployed output is text. STLM's toy also fixes its targets before language-model training and asks whether their externally supplied organization changes what is learned, rather than merely whether a continuous representation can be predicted.

6.7 Latent Visual Reasoning

The recent cluster of latent visual reasoning architectures — LVR [46], Mirage [47], Latent Sketchpad [48], and the latent-space reasoning of Monet [49] — introduces latent visual representations within otherwise text-generating pipelines, commonly as an invoked or interleaved workspace; the degree of compulsory use varies by architecture.

The distinction from our program is the causal role of the intermediate. These systems ask what an interleaved visual workspace contributes to ordinary text generation; STLM asks what remains possible when a substrate representation is compulsory at every lexical position. This is the invoked-versus-constitutive distinction: the substrate is not an auxiliary tool beside the path to speech but the required interface to it. The restriction guarantees mediation, not scene emergence or semantic use. A 2026 causal study reports that removing or randomizing latent visual tokens in existing systems causes little degradation on several spatial-reasoning benchmarks, while arguing that latent supervision may still improve learning even when those tokens are not used at inference [50]. The result does not test STLM, but it reinforces the need to measure causal dependence through compulsory-seam designs or matched ablations, together with permuted or random-substrate controls.

6.8 Pixel-Based Language Models

PIXEL [51] renders text as pixels and processes them with a Vision Transformer. The model operates on images, but they are images of text rather than representations of what the text describes. Text-as-image compression approaches (SEEKER, Glyph, DeepSeek-OCR) [52, 53, 54] apply similarly.

6.9 Autoregressive Text-to-Image and Text-to-Video

Fluid [55], LlamaGen [56], and NextStep-1 [57] use frozen visual encoders in autoregressive next-token prediction, but they predict visual tokens (image patches, codebook indices), driven by text conditioning. Our program inverts the prediction target: we predict text tokens from a substrate built by applying the same kind of frozen encoder to a substrate vocabulary indexed by those text tokens. The combination — per-position autoregressive text generation through a visual, video, or encoded-description vocabulary — is what makes the model a substrate-translated language model rather than a visual generator.

7 Empirical Hypotheses

The program groups its predictions into five hypotheses, each with concrete subpredictions ordered by experimental stage. The unifying prediction is that the substrate produces capability gains in two distinct families: *coherence-specific* gains on tasks where current models fail through coherence loss (long-context drift, character tracking, contradiction detection, narrative-arc maintenance, configurational consistency, schema composition), and *configurational-reasoning* gains on tasks where understanding depends on manipulating relations inherited from physical or schematic source domains (abstract reasoning under perturbation, conceptual analogy, novel-domain transfer). Approximate parity is expected on tasks where neither coherence nor configurational understanding is the bottleneck, such as short-form factual recall, single-sentence completion, and surface fluency. Uniform gains that persist unchanged under permuted or random targets would instead indicate a generic capacity, optimization, or regularization effect.

H1: Operational viability and trade-offs. Phase-zero vocabulary construction should produce stable targets; the SP should learn to predict them from context; and a substrate-only TR trained on the SP's actual outputs should keep the complete system predictively useful without a token-side channel. This feasibility prediction applies to both pixel renderings and frozen text-encoder sequences and is tested by the toy. H1 predicts approximate parity with a parameter-matched direct-token model on standard language benchmarks; the toy provides adverse but confounded evidence (§8.5).

H2: Meaning is causally load-bearing. With substrate shape, capacity, optimization, and readout decodability matched, correctly aligned word-to-substrate targets should outperform fixed permutations and random targets on held-out language modeling and especially on structure-sensitive tasks. Equivalent aligned and permuted conditions would indicate a codebook rather than a contribution from intended meaning. Two ablations refine this claim. First, replacing spatially structured targets with global text-conditioning embeddings of the same dimensionality should reduce performance on spatial benchmarks if explicit geometry matters. Second, adding a direct context-to-logit bypass may improve raw token accuracy but should reduce the effect of ablating or corrupting the substrate. This bypass tests dependence on the seam; only alignment controls test dependence on meaning. Medium comparisons become interpretable after those controls: a pixel advantage on configurational tasks would implicate explicit geometry, whereas comparable aligned effects would implicate structure shared across media.

H3: The substrate is faithful and compositional. The SP's emitted states should preserve the intended upcoming content and structure of the target medium. Raw pixels are directly displayable; image- or video-autoencoder latents can be decoded through their frozen decoder; text-conditioning states such as MiniLM require a separately trained conditioned generator rather than an assumed inverse. Emitted geometry should reflect visual or configurational similarity rather than only lexical co-occurrence, and decoded video should preserve temporal coherence rather than merely per-frame content. Under matched conditions, multiscale targets should improve held-out composition over word-level targets; aligned multiscale targets should outperform a consistent type-preserving permutation; and a consistent permutation should outperform scale-inconsistent permutations if stable factor identity matters across scopes. These orderings test the multiscale intervention, alignment, and cross-scale consistency, not scene emergence. Visually decodable outputs should also depict held-out combinations rather than collapse to familiar exemplars, while H5 tests whether the model uses that structure.

H4: Temporal extent, history, and persistence add task-relevant information. Access to earlier emitted substrate states should improve prediction beyond the current state alone in a matched full-causal-versus-diagonal-attention comparison. The reported TR-C configurations outperform TR-S, but architecture, capacity, and supervision density change together, so they do not isolate history (§8.5). A history gain would show that earlier states carry useful information, not that they are interpreted semantically or form a scene; temporal mismatch and H2's alignment controls remain necessary. At Stage 3, video should outperform still images on action, motion, physical causation, and event sequencing. At Stage 5, persistent state should improve entity, spatial, temporal, and abstract-relation continuity relative to matched nonpersistent substrates.

H5: Structured substrates produce task-specific cognitive effects. After semantic alignment has been validated, structured substrates should improve specifically the tasks whose relevant relations they represent, relative to matched direct-token and permuted-target controls. At toy scale, a negative result weighs against H5 for that configuration, although capacity diverted to substrate prediction limits how far it generalizes. A spatial substrate should improve spatial and physical reasoning; editing persistent state to encode a counterfactual should change the continuation accordingly; and patching or ablating a viewpoint factor should selectively affect perspective-sensitive tasks. Persistent state should improve abstract and relational continuity relative to a matched nonpersistent substrate, while interventions

on schematic factors should selectively affect those predictions and spare matched concrete-object predictions. These claims require controlled behavior and causal interventions, not probe separability or decodability alone.

8 A Stage-1 Toy Implementation: Two Substrate Media

We have implemented stage 1 at toy scale on a single laptop in two substrate variants. Both have the same architectural shape (the SP reads text and emits a per-position substrate; the TR reads only substrate states and emits a token) and the same 1024-token vocabulary (1,021 words drawn from TinyStories plus three special tokens; the lexical words have LLM-authored descriptions). They differ in the *medium* of the intermediate: a *pixel substrate* containing a literal 64×64 RGB image, and a *text-description substrate* containing a fixed-length sequence of frozen text-encoder hidden states over the same description. One set of descriptions is used either as a rendering recipe or as text-encoder input. This paired construction is useful for comparing complete pipelines, but the reported run includes only semantically aligned targets. The toy uses phase 0 for shared vocabulary construction and calls its two post-vocabulary stages *steps*: step 1 trains the SP with §3's phase-two objective in regression form, and step 2 trains TR-S on the SP's actual outputs. Section 3's phase-one canonical pre-training is omitted.

At the 44.8M-token budget, evaluation covers five configurations: the direct-token baseline, pixel and description SP + TR-S, and pixel and description SP + TR-C. The four substrate configurations reuse two frozen SP checkpoints: TR-S and TR-C are alternative readouts over the same medium-specific predictors, while the direct-token baseline is trained separately. The rows therefore compare complete readout configurations, not five independently trained substrate predictors. Section 8.5 gives the full TR-C protocol.

8.1 The Shared Phase-0 Vocabulary

For each of the 1,021 lexical words, a frontier-class language model (used only during phase 0 and then discarded) authors one unique symbolic visual description, typically one or two dense sentences. The three special tokens use fixed sentinel strings in the text condition and fixed solid-color placeholders in the pixel condition. Across all 1,024 encoded text target strings, the inputs occupy a mean of 38.6 content positions (median 37, range 4 to 56) of the $L = 56$ available and span 39 distinct observed lengths; 39 reach the limit, of which 23 exceed it and are truncated. The authoring model supplies the word-to-schema mapping; the downstream diffusion model and text encoder only render or encode the supplied description. *Concrete* tokens get direct depictions; quoting the vocabulary verbatim, "apple" is "A round red apple with a single green leaf attached to its short brown stem, resting on a plain surface," and "dog" is "A golden retriever sits upright on a wooden porch, tail blurred in a wag, mouth open in a panting grin … A well-chewed tennis ball sits at its front paws." *Abstract* and *function* tokens get *image schemas* in the Lakoff-Johnson tradition [8]: "but" is "An ornate iron gate stands open on the left side and firmly shut on the right side … The gate divides two possibilities"; "very" is a thermometer "risen far above the midpoint, nearly to the top"; "now" is "A single bold clock face showing a precise moment … with the word HERE printed in small capitals at the very bottom of the dial."

Distinguishability is an authoring aim rather than an enforced property: the displayed prompt asks for visually specific descriptions and instructs abstract metaphors not to center on vocabulary words, but those constraints do not guarantee global distinctness. "if" is "A signpost at a fork in a dirt road, one arm pointing left labelled with a bold question mark," while "or" is "A forked signpost with two arms of exactly equal length" — two function words using closely related signpost vehicles. More precisely, the lexical token "sign" is itself depicted as a wooden signpost, so the reuse of that object by "if" and "or" is an instance of the vehicle collision described below. The descriptions form a research artifact: 1,021 explicit word-to-description mappings stored as a plain JSON object without schema-type tags. They constitute a falsifiable hypothesis about what cognitive primitives the vocabulary maps onto.

The frontier LLM receives a deliberately minimal prompt:

> *For each word, describe an image representing it. Each description should be a few sentences — visually specific and complete enough to render. Abstract words depicted via metaphor should not center on a vocabulary word.*
>
> *Output: word → "description" (one per line)*
>
> *Vocabulary: [vocabulary list]*

Two separate out-of-vocabulary probes illustrate the intended authoring operation on harder abstractions:

> *emergence* → *"A cluster of iron filings scattered on a dark surface, individually pointing in random directions, but collectively forming a smooth, unmistakable spiral pattern radiating outward from a central magnet hidden beneath the surface."*
>
> *karma* → *"An open ledger lying flat, its left page filled with dense columns of small handwritten figures, its right page bearing a single large number at the bottom that exactly equals their sum, a quill resting across the spine as if just set down."*

These examples are not members of the released vocabulary and are not experimental evidence. They show the intended compression of an abstract concept into a renderable schema: iron filings reuse a familiar physics demonstration, while the ledger frames karma through cumulative accounting. The examples also expose a risk of the method: an authoring model can import culturally specific explanatory conventions, which should be recorded and audited rather than treated as universal.

The minimal prompt leaves composition and metaphor selection largely unconstrained, preserving variation but reducing experimental control. The one substrate-specific constraint — the vehicle-collision rule, a property of this architecture rather than a general rule for description writing — is stated explicitly. The prompt may still admit culturally conventional or lexically identifying depictions (§8.1). The resulting vocabulary should therefore be treated as one authored target set, not as a privileged mapping from words to cognitive primitives.

The vehicle-collision constraint is specific to substrate-based architectures and would not arise in token-LM design. When an abstract word's image-schema requires a concrete object that is itself a vocabulary word, the two canonicals may share the same dominant visual and become hard to distinguish at the substrate's resolution. Avoiding such collisions requires selecting vehicles outside the lexical inventory or explicitly modeling shared visual components. Vocabulary authoring is therefore part of substrate design rather than a mechanical preprocessing step.

Class separability and semantic structure. The authoring procedure favors visually specific canonicals, and the vehicle-collision rule pushes metaphor vehicles away from vocabulary objects. Both choices can help the TR by increasing mutual distinctness, but they can also weaken the semantic neighborhood structure the hypothesis requires. A mutually distinct target set makes a good error-correcting codebook, while the meaningful-substrate claim requires semantically related words to occupy *related* regions. An iron gate for "but" and a forked signpost for "if" are interpretable to a person, but their geometric relation is not an independently justified geometry of conjunction and conditionality; it may be an LLM-designed mnemonic. Four properties should therefore be kept apart: the human interpretability of an individual canonical, the class separability of the set, the semantic neighborhood structure across the set, and compositional factor sharing between canonicals. Only the last two are what the hypothesis needs, and neither is established by inspecting examples. Low-cost diagnostics include pairwise distance distributions, nearest-neighbor listings, correlation with an independent semantic-similarity measure, and within- versus between-schema distances. A vocabulary built from reusable factors (object, color, location, containment, direction, motion) with held-out *combinations* would test compositional sharing far more directly than 1,021 individually authored lexical canonicals.

Target-set effective rank. A further diagnostic is the participation-ratio effective rank of the target matrix's singular spectrum, $(\sum_i \sigma_i)^2 \,/\, \sum_i \sigma_i^2$. Flattened, the pixel canonicals form a $1024 \times 12{,}288$ matrix with values scaled to $[0, 1]$, and the description canonicals a $1024 \times 21{,}504$ matrix. The stored description states are already unit-normalized per position.

Target set	Uncentered	Mean-centered
Pixel canonicals	57.9	193.0
Description canonicals	564.3	769.0

Uncentered, the description set spreads its variance over roughly ten times as many directions as the pixel set; after mean-centering, the ratio falls to about four. All 1024 pixel canonicals are distinct, so their concentration is not duplication. It is consistent with shared image statistics such as a common mean, global brightness or color cast, and broad layout, although the leading singular vectors were not inspected. Effective rank does not measure numerical conditioning, but the description set's less concentrated output code is a nonsemantic candidate explanation for the medium gap in §8.5.

Lexical leakage in the description targets. The authored descriptions were not written under a name-blind constraint, so some contain the target word itself. A case-insensitive whole-word audit of the shipped file finds that **17.9% of the**

1,021 lexical descriptions (183) contain their own target. This raw count does not distinguish label-revealing uses from incidental matches in ordinary prose; both "apple" and "dog" are explicit self-mentions, while short function words such as "a" and "the" can also match incidentally. The channel is therefore real but not universal, and its effective strength cannot be inferred from the aggregate rate alone. For the affected words the frozen encoder receives the target's own name as input, and the encoded canonical is in part a contextualized representation of the token label rather than a purely non-lexical meaning configuration. This does not affect the operational claim, but it weakens the statement that the description targets' organization is specified independently of token identity, plausibly inflates the description medium's advantage over pixels (whose rendered output contains no string), and makes the description pipeline resemble a continuous token-label channel. The required controls are: regenerated *name-free* descriptions (no target string, lemma, inflection, or spelling fragment, with the leakage rate audited before and after filtering); a *name-only* substrate encoding just the target word, to measure how much direct lexical identity explains; and a *masked-name* condition replacing target occurrences with a generic placeholder. For the description medium, the informative comparison is name-only versus name-free versus permuted name-free.

The causal test of meaningfulness. A unique canonical for each word can always act as a codeword, even if its content is meaningful to a person. The first causal alignment experiment therefore keeps the tensors and training procedure fixed while changing only their assignment. Let S_{align} use each word's intended image or description, S_{perm} apply one of several independently sampled fixed permutations of the lexical canonicals while keeping special-token targets fixed, and S_{random} use capacity-matched synthetic targets. Define the alignment effect on an evaluation score M as

$$\Delta_{\text{meaning}} = M(S_{\text{align}}) - M(S_{\text{perm}}).$$

A positive and replicable Δ_{meaning}, especially on compositional or configurational transfer tasks and under matched canonical-decoding accuracy, is evidence that semantic organization is being leveraged. Even then the interpretation must be graded: aligned targets can win because semantically related words receive nearby targets (easing regression under uncertainty), because the target set forms an effective error-correcting output code [11], or because similar classes share statistical strength — real effects, but weaker readings than configurational computation. A positive alignment effect on ordinary language modeling should therefore first be described as evidence for a *semantically organized output geometry*; claims about scene construction or simulation require compositional and intervention tests in addition. Equivalent aligned and permuted performance means that the SP and TR have learned an arbitrary communication code. The present pilot evaluates only S_{align} and therefore cannot estimate Δ_{meaning}.

Even with one fixed canonical per word, alignment affects the learned function, so the permutation control remains informative. The regression analysis of §3.3 shows why: the SP approximates $C^\top p(\cdot \mid x)$, so the geometry of the target set and the assignment of words to points within it both enter the learned function. Alignment can therefore affect interference between words sharing parameters, what the conditional mean looks like under next-token uncertainty (averaging two dog canonicals still yields something dog-like; averaging unrelated canonicals yields a state near neither), how separable those blends are to the TR, and what generalizes. A fixed permutation preserves the target set exactly while changing that assignment, so it remains a valid control in this architecture.

The present toy has limited power to detect this effect: the vocabulary is small, every word recurs often enough that memorizing 1024 mappings is feasible without exploiting shared structure, the canonicals were authored for mutual distinctness rather than for semantic neighborhood structure (§8.1), and aggregate next-token metrics are dominated by function words and `<unk>` where target meaning should matter least. The two outcomes therefore carry unequal evidential weight. A positive Δ_{meaning} replicated across training seeds and independent fixed mappings would be informative. A null would establish a sensitivity limit for this pilot, not that semantic grounding is useless — and in particular would not bear on the multi-prototype, context-selected, or persistent-state variants (§2, §5) in which meaning has substantially more opportunity to become load-bearing. Evaluating compositional minimal pairs instead of aggregate accuracy would concentrate the predicted effect and increase the power of the test.

8.2 Pixel-Substrate Implementation

The pixel-substrate variant uses a literal 64×64 RGB image as the per-position substrate. Each of the 1,021 LLM-authored descriptions is rendered into an image by an open-weights text-to-image diffusion model (SANA-Sprint 1.6B [58] in our implementation); those images plus three fixed special-token placeholders are the 1,024 phase-0 canonicals the SP will learn to emit. Pixels are used directly rather than a CLIP-style global embedding because the architecture's spatial-axis requirement (§9) is preserved: a $3\times64\times64$ tensor has explicit (x, y) structure; a CLIP-512 global vector does not. The pixel SP has around 16M parameters, most in the linear head from its transformer state to the $3\times64\times64 = 12{,}288$ output pixels; TR-S is a small CNN of around 2.5M parameters.

Resolution directly determines the cost–legibility trade-off because the pixel head scales quadratically with image width while the transformer body stays fixed. Figure 2 and Table 1 show that 64×64 retains the illustrative apple's silhouette, color, and stem without making the model mostly a pixel projector. The 64×64 choice is based on this one illustrative word; a canonical-decoding ablation across 32×32 and 64×64 is needed to establish a general threshold.

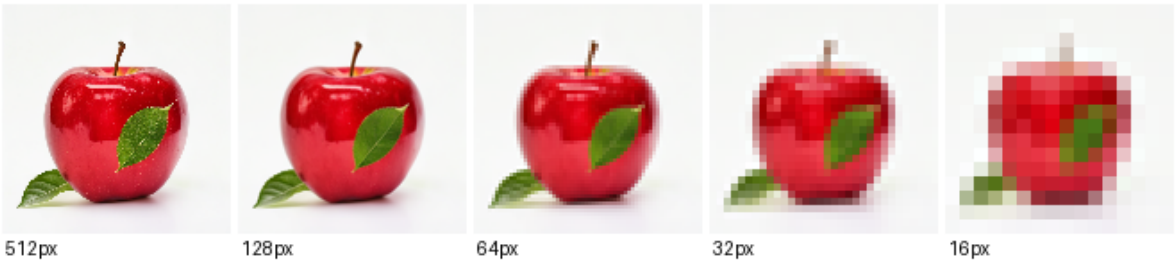


Figure 2: The same canonical ("apple," generated once by Sana-Sprint and downsampled) at five substrate resolutions: 512, 128, 64, 32, 16 pixels per side. The pixel substrate trades off iconic legibility against output-head cost. The toy operates at 64×64, the smallest resolution that retains the apple's distinguishing structure while keeping the SP's output head from dominating the parameter budget (Table 1).

Resolution	Head params	SP total	Head share	Canonicals on disk
16×16	0.3 M	11.4 M	3%	0.8 MB
32×32	1.2 M	12.2 M	10%	3.1 MB
64×64	**4.7 M**	**15.8 M**	**30%**	**12.6 MB**
128×128	18.9 M	29.9 M	63%	50.3 MB
256×256	75.5 M	86.6 M	87%	201 MB
512×512	302 M	313 M	96%	805 MB

Table 1: Parameter and storage scaling of the toy across pixel-substrate resolutions, with the transformer body fixed at 6 layers, $d_{\text{model}} = 384$. "Canonicals on disk" is the full per-word canonical set (1024 images, uint8). The head share is the fraction of the SP's parameters consumed by the pixel-projection head; above 128×128 the model is mostly head. The toy uses 64×64 (bold).

Training has two steps. **Step 1:** train the SP alone — a causal transformer with a pixel-output head — to predict the canonical pixel image of the next token from the running text. Loss is per-pixel MSE. **Step 2:** freeze the SP; sample next-token contexts and project only its final-position pixel output for each context; train TR-S (a small CNN with BatchNorm) with cross-entropy to map those outputs to the correct next token. TR-S thus tunes to the SP's actual output distribution rather than to the idealized canonical distribution — inference-time and training-time inputs match exactly. The toy uses a strictly sequential pipeline, simpler than the joint or lockstep variants §3 entertains.

Phase 0 rendered the 1,021 lexical canonicals at approximately 50 seconds per image via Sana-Sprint on the M4's MPS backend and added three fixed special-token placeholders, totalling approximately 14 hours as a one-time preprocessing step. In the reported run, step 1 completed an epoch-equivalent token budget sampled as random windows with replacement (21,877 optimizer steps, batch 16, context 128) in 48 minutes. Step 2 processed 100,000 sampled next-token contexts (12,500 optimizer steps, batch 8) in 5 minutes 19 seconds. With canonicals already cached, the pixel half of the experiment therefore trained in approximately 53 minutes.

8.3 Text-Description-Substrate Implementation

The text-description substrate uses the same authored descriptions but encodes them rather than rendering them. In the intended design, this medium keeps ordered roles, relations, mechanisms, and consequences available where literal depiction is strained; the toy tests only fixed encoded targets. A frozen text encoder (`all-MiniLM-L6-v2` [59]) maps each description to a fixed-length sequence of per-token hidden states, with $L = 56$ and $d_{\text{enc}} = 384$ in the toy. The full sequence is retained rather than mean-pooled because pooling would collapse word order and argument binding. Padded positions use one fixed unit-norm PAD vector from the encoder's padding-token embedding. Every slot therefore has a supervised target: a contextual token state or the shared PAD direction.

The architectural shape matches the pixel variant. The head differs: rather than projecting to $3{\times}64{\times}64 = 12{,}288$ pixels, the SP broadcasts its hidden state $h(x) \in \mathbb{R}^{d_{\text{model}}}$ across the L description slots, adds a learned per-slot embedding e_j,

and applies one *shared* projection $W \in \mathbb{R}^{d_{\text{enc}} \times d_{\text{model}}}$, so slot j emits

$$g_j(x) \;=\; \frac{W\,\mathrm{LN}(h(x)+e_j)}{\left\|W\,\mathrm{LN}(h(x)+e_j)\right\|}.$$

The slot embeddings make the L outputs distinct; a shared projection alone would emit the same vector at every position. This head costs $d_{\text{model}}d_{\text{enc}} + Ld_{\text{model}} \approx 169$K parameters, rather than the $d_{\text{model}}Ld_{\text{enc}} \approx 8.3$M required by a full projection to the flattened sequence. The description SP is consequently the smaller predictor at 11.3M parameters, compared with 15.8M for pixels. Its loss is cosine distance over the encoder-state sequence rather than per-pixel MSE.

TR-S is a small sequence reader rather than a CNN. It adds a learned class token and positional embeddings, applies a one-layer bidirectional Transformer encoder, and maps the class state to vocabulary logits. It receives no story-context tokens or target-derived attention mask; PAD is represented only by its fixed direction inside the substrate. The description pipeline also avoids the large categorical softmax that an explicit English-token substrate would require.

A second code channel is *length*. Content occupies a mean of 38.6 of the 56 positions, so roughly a third are PAD; because each word has one fixed canonical, the content/PAD boundary exposes description length, and withholding an explicit attention mask from TR-S does not close this channel, since the repeated PAD direction makes the mask inferable. Length is not by itself identifying: the vocabulary spans only 39 distinct lengths, about 5.3 bits against the 10 bits needed to name one of 1024 classes, so it narrows the candidate set rather than determining it. The gap between content-only and PAD-inclusive cosine similarity in the reported run (§8.5) nonetheless shows that PAD prediction is a substantial share of the regression objective.

A third channel that cosine training would otherwise open is closed by construction. Cosine distance is scale-invariant, so a predictor trained under it leaves emitted magnitudes free, and 56 class-dependent norms would be readable by the TR without any deviation in direction. In this implementation the SP applies an explicit unit normalization to every emitted position before the seam, so all magnitudes are one and no norm information crosses it. Implementations that omit that normalization must treat per-position norms as a further leakage channel and probe it directly. Diagnosing this channel requires a length-only classifier, a condition with content states replaced by noise but PAD positions preserved, length-equalized or shift-randomized padding, PAD masked out of both losses, order-shuffled content states, and a mean-pooled baseline. Until those run, the description result should be read as encoded canonical-sequence prediction, not as evidence that sequence order or argument binding is doing work.

Training again proceeds in two steps after phase 0. **Phase 0:** encode the 1,021 authored descriptions and three fixed special-token sentinel strings through the frozen text encoder and cache the per-token hidden states plus a diagnostic attention mask. This deterministic, non-learning step completes in seconds. Padding is materialized as the same fixed PAD vector in every padded slot, so the substrate is fully specified at all L positions. **Step 1:** train the SP to predict the encoded-description sequence of the next token's canonical from the running text. Loss is cosine distance over all L positions; the attention mask is used only for diagnostic metrics that separate content-token error from PAD-slot error. **Step 2:** freeze the SP; sample the same next-token contexts used by the pixel variant and project only its final-position sequence for each context; train TR-S with cross-entropy to map those sequences to the correct next token. TR-S never receives the corpus prefix itself: its only input is the SP's noisy substrate sequence. Evaluation reports frequency baselines and top-k accuracy on SP outputs, since beating random is not an informative bar for a 1024-way Zipfian vocabulary.

Phase 0 completes in seconds to minutes (frozen-encoder forward passes). In the reported run, step 1 completed the same epoch-equivalent token budget sampled as random windows with replacement (87,508 optimizer steps, batch 4, context 128) in 1 hour 49 minutes; step 2 processed the matched 100,000 contexts in 4 minutes 33 seconds. Excluding cached phase-0 construction, the six recorded stages of the two-medium run took a total of 2 hours 47 minutes on the M4 MPS backend.

8.4 Comparison and Discussion

The comparison concerns structured continuous description encodings and structured continuous pixels. The two media share the training corpus and evaluation contexts; within each readout scope, their sampled readout-training windows are also shared. Their SP and TR modules are medium-specific. Two SP differences are quantitative and pull in opposite directions. The pixel SP is the larger model (15.8M parameters, 4.7M of them in the pixel head) against the description SP's 11.3M; but it consumed the identical epoch-equivalent token budget in 21,877 optimizer steps at batch 16, against the description SP's 87,508 steps at batch 4 — four times fewer gradient updates. The reported ordering therefore rests on a capacity advantage for pixels and an update-count advantage for descriptions, neither of which is controlled. A result for one medium measures the complete medium-specific pipeline, not a perfectly isolated causal effect of geometry. Attributing a difference specifically to geometric structure requires stronger controls.

A further limitation applies to both media: the pilot's output heads are single dense projections from one transformer state, trained with coordinatewise losses. For any fixed permutation P of output coordinates, $\|PWh - Pv\|^2 = \|Wh - v\|^2$, so nothing in the SP's architecture or objective distinguishes the true pixel arrangement from an arbitrary fixed shuffling of it; the same holds across the description sequence's positions. The medium's geometry is thus a property of the *targets* and of TR-S's inductive bias (the CNN sees locality; the sequence reader sees order), not something the SP's computation must respect. It is useful to separate a *target-shaped interface* — the output tensor can be displayed as an image or read as a sequence — from a *medium-respecting predictor*, whose computation is constrained by locality, translation, or ordering (a convolutional, latent-grid, or diffusion decoder; a shared sequence decoder). The pilot establishes the former; testing rendering pressure (§1) requires the latter.

8.5 Pilot Experiment and Results

Protocol. The reported run uses seed 42, the first 250,000 TinyStories [60] records, a 1024-token vocabulary, and context length 128. Tokenization produces 44,812,714 tokens, of which 9.0% are `<unk>`. Each SP is trained for an epoch-equivalent token budget using random windows sampled with replacement: 21,877 optimizer steps at batch 16 for the pixel SP, 87,508 steps at batch 4 for the description SP. Each TR-S is then trained against its frozen SP on the same 100,000 sampled contexts with batch size 8 and evaluated on the same 1024 deterministic, evenly spaced contexts from a disjoint 8192-token holdout stream. Even spacing places these contexts roughly eight tokens apart while each carries a 128-token window, so neighboring examples overlap heavily and are not independent; the effective sample size is far below 1024. TR-S receives only the SP's final emitted substrate state. The training pipeline requires a completed SP checkpoint before the readout stage begins, so no TR is fitted to a partially trained predictor. The run is recorded as `data5x_seed42`.

The derived TR-C run `tr_c_data5x_seed42_t128` reuses the frozen pixel and description SP checkpoints from `data5x_seed42`. The TR-S checkpoints are re-evaluated only as comparators; they neither initialize nor participate in TR-C. For each of 100,000 deterministically sampled length-128 windows, shared across media, the frozen SP emits its causal state at every position. A freshly initialized medium-specific adapter, shared across positions, maps each state to 256 dimensions. A two-layer, eight-head causal Transformer reads the complete sequence including the current state and emits vocabulary logits at every position.

Training uses batch size 8 and same-index cross-entropy against all 128 aligned next-token targets, for 12.8 million supervised positions per medium and 12,500 optimizer updates. Source-token IDs and SP hidden states are inputs only to the frozen SP and never to TR-C. Evaluation uses the identical 1024 held-out contexts and reports final-position NLL and top-k metrics, so each TR-C result is evaluated on the same outcomes as its medium's TR-S result. Dataset, token-stream, vocabulary, checkpoint, and metrics identities are pinned in the accompanying reproduction record.

Table 2 also reports an earlier run at one fifth of this data budget (`paper_seed42_e1_n2-100k`, 8,826,412 tokens), which uses the same protocol, seed, and canonicals with only the corpus slice changed. Its holdout stream is the tail of a shorter slice and therefore a different text with different token frequencies, so the two runs are compared by margin over their own baselines rather than by raw accuracy.

System	SP train tokens	Eval. NLL	Top-1	Top-5	Top-10
Context-free frequency baseline	44.8M	–	8.1%	27.2%	39.5%
Direct-token LM (backbone-matched)	44.8M	**2.208**	**44.7%**	**74.6%**	**85.3%**
Description SP + TR-C	44.8M	2.616	41.9%	68.6%	78.0%
Pixel SP + TR-C	44.8M	2.954	37.8%	61.6%	72.2%
Description SP + TR-S	44.8M	3.541	35.0%	54.2%	62.8%
Pixel SP + TR-S	44.8M	4.283	27.8%	44.3%	52.9%
Context-free frequency baseline	8.8M	–	11.0%	27.4%	39.7%
Description SP + TR-S	8.8M	3.998	29.7%	47.9%	56.1%
Pixel SP + TR-S	8.8M	4.466	22.0%	40.1%	50.3%

Table 2: Next-token results on 1024 held-out contexts. Within each budget all rows use the same holdout. The 44.8M-token block is the five-configuration comparison: TR-C reuses the frozen SPs but is freshly trained with a different architecture and 12.8M aligned labels per medium, versus 100,000 final-position labels for TR-S, so TR-C–TR-S differences do not isolate history. Across budgets the holdout differs, which is why the frequency baseline moves. All percentage-point differences discussed below are computed from the unrounded accuracies. The context-free baseline always predicts the k most common holdout targets; it is an oracle test-marginal reference, not a deployable model or a mathematical floor. Each trained condition is one seed; no confidence intervals are claimed.

Feasibility result. All four substrate configurations substantially exceed the context-free frequency baseline at 44.8M tokens. TR-S reaches 35.0% top-1 for descriptions and 27.8% for pixels, establishing that the current SP output alone carries context-dependent next-token information. TR-C reaches 41.9% and 37.8%, improvements of 6.9 and 10.0 points over the corresponding TR-S conditions on the identical holdout. Because neither readout has source-token access, the emitted state or trace is sufficient for a separate reader to recover useful predictive signal. The gain does not isolate temporal memory, and none of these results shows that the SP exploited the intended meaning of the targets; fixed canonicals may function as learnable codewords.

Direct-token baseline. A backbone-matched direct-token language model — matching the pixel SP's causal transformer body, corpus slice, seed, batch size, optimizer, learning rate, and 21,877 optimizer steps, but with an ordinary token-softmax head and no seam or TR — reaches 44.7% top-1 and 2.208 next-token NLL (perplexity 9.1) on the same 1024 holdout contexts. It shares the transformer body, corpus slice, and seed with the description SP, whose batch size and update count differ as reported above. The baseline is recorded as `baseline5x_seed42`. Against TR-S it leads descriptions by 9.8 top-1 points and pixels by 16.9; against TR-C the gaps narrow to 2.8 and 6.9 points. Perplexity is 9.1 for the direct model, 13.7 and 19.2 for description and pixel TR-C, and 34.5 and 72.5 for TR-S. The direct system has 11.48M parameters; TR-C adds 2.66M parameters to the description SP and 1.93M to the pixel SP and receives 12.8M position labels per medium. Total FLOPs were not measured. These numbers rank the reported configurations but do not estimate a causal cost for the seam. The baseline does not test context-dependent scene targets, so its lead weighs against H1's parity subprediction only for the present implementation.

The run using a five-times larger corpus slice records higher values on every reported metric, with code, seed, and canonicals held fixed. Top-1 margins over each run's own frequency baseline are 26.9 rather than 18.7 points for descriptions and 19.7 rather than 10.9 for pixels; TR-S cross-entropy is 3.541 rather than 3.998 and 4.283 rather than 4.466. The pixel SP's evaluation MSE is 17% lower (0.0483 rather than 0.0582), and the description SP's content-position cosine similarity is 0.417 rather than 0.379. The larger run sees roughly 2.83 tokens per SP parameter for pixels and 3.98 for descriptions (2.45 and 3.26 if TR-S is counted), far below the roughly twenty-tokens-per-parameter regime reported for much larger conventional transformers [61]; whether that heuristic transfers here is unknown. Both continuations remain degenerate, and further training could improve, plateau, or degrade generalization.

This pattern is consistent with the smaller run being budget-limited, but the two-point comparison is not a causal estimate or a clean learning curve. Each run's holdout is the tail of its own corpus slice, so the evaluations differ in text and intrinsic difficulty; comparing margins over each run's own baseline mitigates but does not control for that. A proper budget comparison evaluates both checkpoints on identical frozen contexts, which for these runs means using the larger run's holdout: the smaller run's holdout lies inside the larger run's training stream, so the reverse direction would score a model on text it was trained on.

Across media, the text-description substrate is stronger at every reported cutoff and with both readout scopes. Under TR-S it exceeds pixels by 7.1 top-1 points; under TR-C the gap is 4.1 points. The SP regression diagnostics are not directly comparable across media: the description SP reaches 0.417 content-position cosine similarity (0.482 including PAD positions), while the pixel SP reaches MSE 0.0483. The evidence therefore favors the description pipeline as an information channel in this particular toy. Without each medium's permuted and random controls, the difference cannot be attributed to ordered semantic structure, spatial geometry, or meaningfulness generally. The pilot trains pixels in a fixed-canonical MSE regime expected to be unfavorable to crisp conditional images. The pixel medium's hypothesized advantage — rendering pressure that may favor scene assembly (§1) — arises, if at all, under generative, context-conditioned objectives at scale, a regime this pilot does not reach. The ordering therefore applies only to these two small pipelines under fixed-canonical regression; the space of candidate media remains open (video, audio, and hybrid substrates; §9).

Qualitative output. From the prompt "the dog ran into the yard," the text variant continues "it was and it was and it was and it was it was and," while the detokenized pixel continuation gives "and was the and and." Both strips also contain `<unk>` predictions that the detokenizer suppresses. These labels are the expected degenerate predictions rather than anything substrate-specific: `<unk>` is the corpus's single most frequent target (9.0% of tokens) and "and" is among the top function words, and an undertrained model facing next-token uncertainty falls into exactly this high-frequency loop in either medium.

The pixel strip changes qualitatively in the larger-corpus run, which narrows a confound the smaller run left open. At 8.8M tokens every emitted state was the same gray-brown average, consistent both with the regression-to-the-mean effect §3 predicts for fixed-target training and with simple undertraining. At 44.8M tokens the early positions are visibly differentiated and partly recognizable: Figure 3 contains a dog-like state, a face-like state, and a distinct red vertical element, before later positions decay into the earlier uniform blur. This pattern is consistent with undertraining contributing to the smaller-run collapse, but the runs use different holdout text and therefore do not identify the cause.

The residual decay is also consistent with degenerate generation contributing: once the emitted text enters the and/<unk> loop, the predictor conditions on its own repetitive prefix, so its predicted target distribution can also repeat.

Two cautions attach to this improvement. First, legibility is not yet fidelity: the decoded tokens do not match what a human reads in the states, and a state that strongly resembles the vocabulary's authored dog canonical is decoded as <unk>. Because this is a free-running continuation, there is no ground-truth next token at that position; visual resemblance also does not establish that the SP emitted the correct canonical or that TR-S alone is at fault. This sample establishes only that human-recognizable dog-like structure and machine decodability can diverge. The mismatch is consistent with TR-S reading artifact dimensions rather than gross content, but equally with an off-canonical dog, with insufficient readout invariance to a narrow output distribution, or with an ordinary class-boundary error. Separating these requires the aggregate intervention of the channel-versus-manifold readout comparison (§9), not inspection of individual frames. Second, the qualitative sample remains a failure analysis, not evidence of coherent generation.



Figure 3: The first nine pixel-substrate states emitted after the prompt "the dog ran into the yard" in the 44.8M-token run, labelled with TR-S's decoded tokens. Unlike the smaller run, in which every state was the same gray-brown average, several early states are differentiated and partly recognizable — a dog-like state in the sixth position, a face-like state in the second, a red vertical element in the third — before later positions decay toward uniform blur as the generated text enters its and/<unk> loop. The decoded tokens do not correspond to what a human reads in the states: the dog-like state is decoded as <unk>. This is an unaltered crop of the run's substrate-state strip.

The text-description substrate is less directly inspectable than an explicit English-token sequence but more diagnostically accessible than an unstructured latent: nearest-neighbor lookup can be applied per encoder position or over whole canonical sequences to approximate what region of description space the SP has entered. This is a research aid, not the substrate's purpose. The pixel variant is directly displayable, but displayability is not the same as interpretability, as Figure 3 shows.

8.6 Limitations

The toy deliberately uses one fixed canonical per word and trains both readouts on actual SP outputs. This simplifies fidelity analysis while enabling fixed-code and artifact shortcuts; raw pixels and frozen description states likewise support medium-specific probes but do not capture higher-dimensional reconstruction latents.

Scope of the evidence. The pilot supports only narrow causal and comparative conclusions. It contains no accumulating scene, uses one seed per condition and two corpus budgets with different holdout text, maps 9.0% of tokens to <unk>, and evaluates on heavily overlapping contexts (§8.5). It therefore supplies no uncertainty estimates or population-level comparison.

The design also lacks a capacity- and compute-matched direct-token baseline, permuted or random-target control, earlier-state intervention, and targeted spatial or coherence benchmark; its media and readouts differ in dimensions, losses, capacity, updates, and supervision density. Future runs should split by story, sample non-overlapping contexts, report accuracy excluding <unk> and by token-frequency bin, and prefer a tokenizer without a large out-of-vocabulary class. The reported numbers are single-seed, plausibly budget-sensitive configuration estimates rather than lower bounds.

9 Future Work

Immediate validation. Three near-term steps answer different questions before video or frontier scale. First, isolate history use in the existing pilot with the matched H4 ablation: compare full lower-triangular attention with diagonal current-state-only attention while holding the frozen SP emissions, adapter, Transformer, parameter count, initialization seed, data order, training and evaluation windows, aligned position losses, optimizer, and update count fixed. Multiple paired readout seeds should share their initial weights within each pair and report paired per-example estimates with uncertainty intervals. A paired gain would show that earlier emitted states contain usable predictive information; it would not establish semantic use, accumulation, or scene formation. Second, estimate Δ_{meaning} using fixed permutations under

the present design (§8.1). Third, test Stage 2 by replacing the atomic mapping with random and then context-selected prototype instances; the same alignment controls then become sharper tests. Stage 4's separate multiscale experiment follows below.

For any alignment experiment, the minimum design begins with a matched direct-token LM. Within each substrate medium, aligned, permuted, and random targets must use identical training data, sampled contexts, holdout, architecture, and optimizer-update count, so that no condition receives more gradient steps than another (§8.4). Readout decodability must also be matched, so that an easy codebook does not masquerade as a semantic effect.

Paired aligned and permuted runs should share initialization and data order, and they should be crossed with multiple independently sampled fixed permutations that keep special-token targets fixed. At least three training seeds, per-example predictions, and uncertainty intervals are required. Evaluation should pair ordinary language loss with a controlled compositional or configurational benchmark. For descriptions, it should also include the name-only, masked-name, and name-free leakage conditions of §8.1 together with the padding and length ablations of §8.3.

Controlled multiscale experiment. A small factorial microworld can test cross-scale reuse more directly than open-ended TinyStories. Its grammar should generate scenes from stable object, color, location, relation, and motion primitives, with the same components reused in word-, phrase-, sentence-, and short-discourse targets. Complete combinations should be withheld during training while every constituent appears in other combinations. The comparison comprises a matched direct-token LM, a word-level substrate, a multiscale aligned substrate, a multiscale substrate using one fixed type-preserving permutation consistently at every scale, and a multiscale substrate using independent within-type permutations at each scale. Objects must permute among objects, colors among colors, relations among relations, and motions among motions; the split must be balanced or closed under these permutations so the target multiset and marginal frequencies remain matched. All substrate conditions should match target dimensionality, readout capacity, supervised positions, data, optimizer updates, and decodability as closely as possible; any remaining mismatch should be reported as a whole-configuration comparison rather than a pure effect of composition.

The consistent global permutation preserves target-set geometry and cross-scale reuse while breaking the intended lexical alignment. Multiscale aligned versus consistently permuted therefore estimates the contribution of that alignment. Consistently permuted versus scale-inconsistent tests whether preserving factor identity across levels contributes beyond within-level structure. Multiscale aligned versus word-level aligned estimates the value of the complete multiscale intervention, not factor reuse alone unless supervision and representational capacity are fully matched. Evaluation should report held-out-combination prediction, attribute–object binding, entity persistence, spatial and motion consistency, and responses to controlled edits of color, position, relation, and motion. A matched direct-token LM measures utility. No performance comparison alone licenses a claim of scene emergence: the latent claim also requires state-level evidence for coexisting bound constituents, factor-specific causal edits or patches, and retention or update across positions; the explicit claim further requires direct intervention on a persistent scene.

Each substrate condition should be evaluated under two readout-training regimes: a *channel readout* trained on the SP's outputs as in the pilot, and a *manifold readout* trained on canonical content with meaningful within-concept variation and frozen before evaluation. The latter needs several instances per word where available, not merely perceptual jitter, which would leave it a memorized lookup rather than a semantic reader. The channel readout is expected to score higher because it sees the SP's actual output distribution and may exploit artifacts. The informative quantity is the gap rather than either level: a small gap means the SP's emissions sit near the authored canonicals, while a large one means the decodable information lives in artifacts rather than in the intended content.

Together these comparisons estimate three distinct quantities: alignment, from aligned versus permuted and random targets under matched adaptive readouts; fidelity, from the frozen manifold readout; and utility, from the matched direct-token LM. Because a permutation preserves target-set geometry and changes only the word-to-target assignment, any one-way decoder-adaptation advantage remains available in both aligned and permuted conditions, making Δ_{meaning} under matched adaptive readouts the primary alignment estimator. The frozen reader cannot by itself establish that meaning is load-bearing, and an aligned advantage visible only under the adaptive reader would indicate decoder adaptation or artifact exploitation rather than fidelity. A conditional or generative pixel objective with a medium-respecting decoder (§8.4) should then test whether crisp, context-specific scenes create stronger rendering pressure than the present fixed-canonical regression; for pixels, decoding robustness under compression, blur, quantization, and added noise separates recognizable content from machine-readable residue.

The five-stage path should follow the observed failure rather than a fixed schedule. Failure to represent situated lexical meaning motivates varied prototypes; weak temporal transformation motivates video; weak cross-scale composition motivates span-level or multiscale targets; and weak persistence, inspectability, or manipulability motivates an explicit revisable scene. Positive evidence from a simpler condition can make later machinery unnecessary for some capabilities.

A Stage-5 implementation would still require interventions against recurrent private code; frontier-scale resources should follow controlled evidence rather than feasibility alone.

Parallel STLM–LM selection. At any stage, a deployment hybrid could run an STLM beside a direct-token language model on the same generated prefix. At each position, both paths would predict a next-token distribution and offer a candidate token; a learned or calibrated selector would choose one, append it to the direct model's token history, feed it to the STLM predictor's scene-update path, and repeat. The selected token would never enter the STLM readout directly. Learned token-level deferral between language models has direct precedent in Co-LLM, while LLM-Blender ranks and fuses complete language-model outputs [28, 29]. The additional issue here is architectural: one branch produces language only through a substrate, and its persistent state must remain synchronized even when the direct branch wins. Stage 5 may be the most interesting case because one candidate is then informed by a single scene that persists across the complete book or document; whenever the direct-token candidate wins, the selected token must be assimilated into that scene before the next prediction so the paths do not diverge. Compulsory mediation remains true inside the STLM branch, but the combined system is not itself an STLM under the present definition because final language can bypass the substrate. Evaluation must compare it with an equal-compute ensemble of two direct-token models, report branch-selection rates, and test whether controlled substrate edits change selection appropriately. Token-level routing should also be compared with span-level routing because locally preferred tokens need not form the best coherent continuation.

Tokenization and semantic granularity at scale. The pilot uses a 1024-token vocabulary whose 1,021 lexical entries are whole words, but production models operate on subword fragments, punctuation, whitespace variants, and bytes, many of which denote nothing on their own; there is no canonical image of a suffix or a formatting token. This challenges one-substrate-object-per-token directly, and it is a design problem rather than an implementation detail. Candidate resolutions include word-level substrate generation with a subword or character spelling decoder beneath it; emitting substrate states only at semantic-unit boundaries; predicting multi-token lexical units or concepts, as in next-concept prediction [22]; or letting the TR map one substrate state to a variable-length token sequence. These choices also set the substrate's compositional scope. Comparing lexical, phrase-, and sentence-level targets built from the same factors can therefore test whether a larger semantic unit supplies useful joint structure rather than merely repairing tokenization. Which design preserves the no-bypass property without reintroducing a token-statistics pathway remains open.

Encoder family and spatial structure. For spatial-medium variants, the substrate must expose an explicit spatial axis if the architecture is to support claims about spatial structure. A single global vector — the output of contrastive encoders such as CLIP, SigLIP, or EVA-CLIP — has no (x, y) coordinates internal to it. Contrastive training is caption-discriminative, and captions usually identify content rather than specifying its precise layout, so such encoders plausibly attenuate fine spatial detail their objective did not require. *Reconstruction* latents — the spatially structured grids inside FLUX, Stable Diffusion 3, and masked autoencoders — preserve an axis because they must decode back to pixels. Raw pixels preserve it natively at higher dimensional cost. Even so, an explicit axis is only a precondition: §8.4 shows that it does not force the predictor's computation to respect geometry. Spatial claims therefore require substrates with explicit indexing; global contrastive vectors remain useful comparison media but do not meet that criterion.

At each stage, the borrowed visual and video encoders are optimized for visual or video generation, not for downstream language modeling. A natural extension is to train encoders specifically for the architecture, with compression profiles tuned for language modeling while remaining decodable for the fidelity probes under H3.

The trajectory also generalizes beyond vision. An audio substrate (one or several sound instances per word at Stages 1–2, short audio sequences at Stage 3, span-level sound at Stage 4, and a persistent audio scene at Stage 5) would test whether prediction can be mediated by sound-structured working states. An embodied substrate (text-to-action) would test the same hypothesis with motor-plan structure. The vision-and-video case is the one the necessary encoders most clearly support today; other modalities follow as their encoders mature.

10 Conclusion

The long-term STLM hypothesis is a model that builds a world as it writes: a persistent configurational scene that accumulates interacting physical and schematic structure and is the sole lexical interface. The five-stage path makes the escalation explicit: one fixed substrate per word; varied context-sensitive instances; short lexical clips; substrates spanning several words; and finally one continuously revised scene. Simpler stages may already bias ordinary contextual integration toward a latent scene-like organization, and successful earlier mechanisms can make later ones unnecessary

for some capabilities. This paper implements only the Stage-1 prerequisite: a compulsory no-bypass interface with independent per-position image and description targets.

Synchronized scene–text trajectories can supervise the continuous endpoint. Because a scene underdetermines wording, evaluation should report language-model likelihood separately from complementary tests of semantic adequacy, persistent consistency, and causal response to scene interventions.

Choosing and constructing those trajectories is upstream of the no-bypass claim. STLM makes a structured substrate obligatory and asks whether its intended organization is causally useful; document-to-world compilation asks which synthetic explanatory representations to supply and how to scale them across a corpus.

In the reported TinyStories comparison, both media transmit useful next-token signal through both readout scopes, but the direct-token model remains best. The causal readouts narrow the gap without isolating memory because architecture, capacity, and supervision density change together. The pixel strip also separates human-legible content from machine decoding. These results establish the compulsory seam, not load-bearing meaning or scene emergence.

The next tests separate four questions. A matched full-history versus current-state-only readout tests whether earlier substrate states are useful. Intended mappings versus fixed permutations test whether target alignment matters. Varied lexical prototypes test whether Stage 2 carries within-word meaning rather than one fixed identity. A controlled multiscale microworld tests whether stable objects, attributes, relations, motions, and schemas transfer to held-out combinations. None alone demonstrates scene emergence; causal edits must ultimately show that intended structure drives corresponding behavioral changes.

The stages are escalation options rather than requirements. Varied prototypes are warranted when fixed lexical targets collapse situated meaning, video when still media fail to carry transformations, multiscale targets when lexical targets fail to compose, and explicit persistence when latent integration fails to maintain, expose, or manipulate a coherent state. The program succeeds only if some version becomes more than an expensive code channel: its reusable meaningful organization must be causally responsible for behavior.

Code and Artifact Availability

Code, paper source, vocabulary metadata, immutable run manifests, compact metrics, and per-example TR-C predictions accompany the paper at `https://github.com/emergent-wisdom/substrate-language-modeling`. Canonical target tensors and images, together with the exact eleven checkpoints across the four reported run IDs, are hosted in the companion artifact repository at `https://huggingface.co/emergent-wisdom/substrate-language-modeling`. The tracked release records and the artifact repository's `SHA256SUMS` manifest identify released files by path, byte size, and SHA-256 digest.

The reported runs use the TinyStories revision `f54c09fd2331`. Their manifests bind the dataset and encoder revisions, token streams, vocabulary, canonical tensors, training settings, metrics, and checkpoints by hash. Exact saved-checkpoint evaluation reproduces the reported TR-S measurements at full precision; the TR-C prediction files reproduce top-k counts exactly and NLL within 10^{-7}, the tolerance needed after serializing per-example float32 log probabilities. Fresh training from frozen inputs is pinned and seeded, although Apple MPS does not guarantee bitwise-identical training across hardware or software versions. The two repositories therefore support both external use of the exact checkpoint bytes and independent training replication from authenticated inputs.

References

[1] B. Z. Li, M. Nye, and J. Andreas. Implicit Representations of Meaning in Neural Language Models. *Proceedings of ACL-IJCNLP*, pp. 1813–1827, 2021.

[2] K. Li, A. K. Hopkins, D. Bau, F. Viégas, H. Pfister, and M. Wattenberg. Emergent World Representations: Exploring a Sequence Model Trained on a Synthetic Task. *ICLR*, 2023.

[3] N. Kim and S. Schuster. Entity Tracking in Language Models. *ACL*, pp. 3835–3855, 2023.

[4] J. Feng and J. Steinhardt. How do Language Models Bind Entities in Context? *ICLR*, 2024.

[5] W. Gurnee and M. Tegmark. Language Models Represent Space and Time. *ICLR*, 2024.

[6] L. W. Barsalou. Perceptual symbol systems. *Behavioral and Brain Sciences*, 1999.

[7] A. Paivio. *Mental Representations: A Dual Coding Approach*. Oxford University Press, 1986.

[8] G. Lakoff and M. Johnson. *Metaphors We Live By*. University of Chicago Press, 1980.

[9] H. Westerberg. Entangled Alignment: When Safety Is the Substrate. Zenodo, 2026. doi:10.5281/zenodo.22073296.

[10] H. Westerberg. Imagining the Corpus: Joint Prediction of Language and Explanatory Worlds. Unpublished manuscript, 2026.

[11] J. O'Neill and D. Bollegala. Error-Correcting Neural Sequence Prediction. *arXiv:1901.07002*, 2019.

[12] A. Blattmann et al. Stable Video Diffusion: Scaling Latent Video Diffusion Models to Large Datasets. *arXiv:2311.15127*, 2023.

[13] OpenAI. Video Generation Models as World Simulators (Sora technical report), 2024.

[14] P. W. Koh et al. Concept Bottleneck Models. *ICML*, 2020.

[15] C.-E. Sun, T. Oikarinen, and T.-W. Weng. Crafting Large Language Models for Enhanced Interpretability. *ICML Mechanistic Interpretability Workshop / arXiv:2407.04307*, 2024.

[16] A. A. Ismail et al. Concept Bottleneck Language Models for Protein Design. *arXiv:2411.06090*, 2024.

[17] M. Havasi, S. Parbhoo, and F. Doshi-Velez. Addressing Leakage in Concept Bottleneck Models. *NeurIPS*, 2022.

[18] LCM Team et al. Large Concept Models: Language Modeling in a Sentence Representation Space. *arXiv:2412.08821*, 2024.

[19] C. Shao et al. Continuous Autoregressive Language Models. *arXiv:2510.27688*, 2025.

[20] O. Naparstek. Projected Autoregression: Autoregressive Language Generation in Continuous State Space. *arXiv:2601.04854*, 2026.

[21] S. Kumar and Y. Tsvetkov. Von Mises–Fisher Loss for Training Sequence to Sequence Models with Continuous Outputs. *ICLR*, 2019.

[22] Y. Liu et al. Next Concept Prediction in Discrete Latent Space Leads to Stronger Language Models. *arXiv:2602.08984*, 2026.

[23] H. Tan and M. Bansal. Vokenization: Improving Language Understanding with Contextualized, Visual-Grounded Supervision. *EMNLP*, 2020.

[24] W. Wang et al. Visually-Augmented Language Modeling. *ICLR*, 2023.

[25] W. Zhu et al. Visualize Before You Write: Imagination-Guided Open-Ended Text Generation. *Findings of EACL*, 2023.

[26] T. Tang et al. Learning to Imagine: Visually-Augmented Natural Language Generation. *ACL*, 2023.

[27] P. Ontalvilla, A. Ormazabal, and G. Azkune. Improving the Efficiency of Visually Augmented Language Models. *COLING*, 2025.

[28] S. Z. Shen, H. Lang, B. Wang, Y. Kim, and D. Sontag. Learning to Decode Collaboratively with Multiple Language Models. *ACL*, pp. 12974–12990, 2024.

[29] D. Jiang, X. Ren, and B. Y. Lin. LLM-Blender: Ensembling Large Language Models with Pairwise Ranking and Generative Fusion. *ACL*, pp. 14165–14178, 2023.

[30] J. R. Meehan. TALE-SPIN, An Interactive Program that Writes Stories. *Proceedings of the 5th International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 91–98, 1977.

[31] J. Guan, Z. Yang, R. Zhang, Z. Hu, and M. Huang. Generating Coherent Narratives by Learning Dynamic and Discrete Entity States with a Contrastive Framework. *Proceedings of the AAAI Conference on Artificial Intelligence*, 37(11):12836–12844, 2023.

[32] P. Wang, J. Zamora, J. Liu, F. Ilievski, M. Chen, and X. Ren. Contextualized Scene Imagination for Generative Commonsense Reasoning. *International Conference on Learning Representations (ICLR)*, 2022.

[33] B. Mohapatra, G. Duca, L. Romary, and J. Cassell. Using Machine Mental Imagery for Representing Common Ground in Situated Dialogue. *arXiv:2604.21144*, 2026.

[34] Y. Wang et al. Orca: The World is in Your Mind. *arXiv:2606.30534*, 2026.

[35] S. Hao, S. Sukhbaatar, D. Su, X. Li, Z. Hu, J. Weston, and Y. Tian. Training Large Language Models to Reason in a Continuous Latent Space. *COLM*, 2025.

[36] C. Saharia et al. Photorealistic Text-to-Image Diffusion Models with Deep Language Understanding (Imagen). *NeurIPS*, 2022.

[37] R. Rombach et al. High-Resolution Image Synthesis with Latent Diffusion Models. *CVPR*, 2022.

[38] Black Forest Labs. FLUX.1. Open-weights text-to-image model release, 2024.

[39] Y. LeCun. A Path Towards Autonomous Machine Intelligence. *OpenReview / position paper*, 2022.

[40] M. Assran et al. Self-Supervised Learning from Images with a Joint-Embedding Predictive Architecture. *arXiv:2301.08243*, 2023.

[41] A. Bardes et al. Revisiting Feature Prediction for Learning Visual Representations from Video. *arXiv:2404.08471*, 2024.

[42] M. Assran et al. V-JEPA 2: Self-Supervised Video Models Enable Understanding, Prediction and Planning. *arXiv:2506.09985*, 2025.

[43] H. Huang, Y. LeCun, and R. Balestriero. LLM-JEPA: Large Language Models Meet Joint Embedding Predictive Architectures. *arXiv:2509.14252*, 2025.

[44] D. Chen et al. VL-JEPA: Joint Embedding Predictive Architecture for Vision-Language. *arXiv:2512.10942*, 2025.

[45] S. Wan et al. JEPA-T: Joint-Embedding Predictive Architecture with Text Fusion for Image Generation. *arXiv:2510.00974*, 2025.

[46] B. Li et al. Latent Visual Reasoning. *arXiv:2509.24251*, 2025.

[47] Z. Yang et al. Machine Mental Imagery: Empower Multimodal Reasoning with Latent Visual Tokens. *arXiv:2506.17218*, 2025.

[48] H. Wang et al. Latent Sketchpad: Sketching Visual Thoughts to Elicit Multimodal Reasoning in MLLMs. *arXiv:2510.24514*, 2025.

[49] Q. Wang, Y. Shi, Y. Wang, Y. Zhang, P. Wan, K. Gai, X. Ying, and Y. Wang. Monet: Reasoning in Latent Visual Space Beyond Images and Language. *arXiv:2511.21395*, 2025.

[50] D. Zhu et al. Leveraging Latent Visual Reasoning in Silence. *arXiv:2605.18641*, 2026.

[51] P. Rust et al. Language Modelling with Pixels. *ICLR*, 2023.

[52] Y. Lu et al. From Text to Pixel: Advancing Long-Context Understanding in MLLMs (SEEKER). *arXiv:2405.14213*, 2024.

[53] J. Cheng et al. Glyph: Scaling Context Windows via Visual-Text Compression. *arXiv:2510.17800*, 2025.

[54] H. Wei, Y. Sun, and Y. Li. DeepSeek-OCR: Contexts Optical Compression. *arXiv:2510.18234*, 2025.

[55] L. Fan et al. Fluid: Scaling Autoregressive Text-to-image Generative Models with Continuous Tokens. *arXiv:2410.13863*, 2024.

[56] P. Sun et al. LlamaGen: Autoregressive Model Beats Diffusion: Llama for Scalable Image Generation. *arXiv:2406.06525*, 2024.

[57] StepFun NextStep Team. NextStep-1: Toward Autoregressive Image Generation with Continuous Tokens at Scale. *arXiv:2508.10711*, 2025.

[58] J. Chen et al. SANA-Sprint: One-Step Diffusion with Continuous-Time Consistency Distillation. *ICCV / arXiv:2503.09641*, 2025. Checkpoint `Efficient-Large-Model/Sana_Sprint_1.6B_1024px`.

[59] N. Reimers and I. Gurevych. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. *EMNLP*, 2019. Checkpoint `sentence-transformers/all-MiniLM-L6-v2`.

[60] R. Eldan and Y. Li. TinyStories: How Small Can Language Models Be and Still Speak Coherent English? *arXiv:2305.07759*, 2023.

[61] J. Hoffmann et al. Training Compute-Optimal Large Language Models. *arXiv:2203.15556*, 2022.